
# **Polars User Guide**

None

# Table of contents

1. Polars	5
2. Blazingly Fast DataFrame Library	
2.1 About this guide	5

# 2.2 Performance



2.3	3 Example	6
2.4	4 Sponsors	7
2.5	5 Community	7
2.6	5 Contribute	7
2.7	7 License	7
3. (	Setting Started	8
3.1	Introduction	8
3.2	2 Installation	9
3.3	Series & DataFrames	10
3.4	Reading & Writing	16
3.5	5 Expressions	19
3.6	5 Combining DataFrames	27
4. L	Jser Guide	29
4.1	Introduction	29
4.2	2 Installation	30
4.3	3 Concepts	35
4.4	1 Expressions	50
4.5	5 Transformations	107
4.6	6 Lazy API	132
4.7	7 IO	142
4.8	3 SQL	154
4.9	) Migrating	163
4.1	0 Misc	169

# Polars



# 2. Blazingly Fast DataFrame Library

docs passing Build and test failing crates.io v0.31.1 pypi v0.18.7 npm v0.8.0

DOI 10.5281/zenodo.7697217

Polars is a highly performant DataFrame library for manipulating structured data. The core is written in Rust, but the library is available in Python, Rust & NodeJS. Its key features are:

- Fast: Polars is written from the ground up, designed close to the machine and without external dependencies.
- I/O: First class support for all common data storage layers: local, cloud storage & databases.
- Easy to use: Write your queries the way they were intended. Polars, internally, will determine the most efficient way to execute using its query optimizer.
- Out of Core: Polars supports out of core data transformation with its streaming API. Allowing you to process your results without requiring all your data to be in memory at the same time
- Parallel: Polars fully utilises the power of your machine by dividing the workload among the available CPU cores without any additional configuration.
- Vectorized Query Engine: Polars uses Apache Arrow, a columnar data format, to process your queries in a vectorized manner. It uses SIMD to optimize CPU usage.

# 2.1 About this guide

The Polars user guide is intended to live alongside the API documentation. Its purpose is to explain (new) users how to use Polars and to provide meaningful examples. The guide is split into two parts:

- Getting Started: A 10 minute helicopter view of the library and its primary function.
- User Guide: A detailed explanation of how the library is setup and how to use it most effectively.

If you are looking for details on a specific level / object, it is probably best to go the API documentation: Python | NodeJS | Rust.

# 2.2 Performance 🚀 🚀

Polars is very fast, and in fact is one of the best performing solutions available. See the results in h2oai's db-benchmark, revived by the DuckDB project.

Polars TPCH Benchmark results are now available on the official website.

# 2.3 Example



API scan\_csv · API filter · API groupby · API collect

```
import polars as pl

q = (
    pl.scan_csv("docs/src/data/iris.csv")
    .filter(pl.col("sepal_length") > 5)
    .groupby("species")
    .agg(pl.all().sum())
)

df = q.collect()
```

API LazyCsvReader API filter API groupby API collect • 📮 Available on feature streaming • 📮 Available on feature csv

```
use polars::prelude::*;
let q = LazyCsvReader::new("docs/src/data/iris.csv")
    .has_header(true)
    .finish()?
    .filter(col("sepal_length").gt(lit(5)))
    .groupby(vec![col("species")])
    .agg([col("*").sum()]);
let df = q.collect();
```


API scanCSV API filter API groupBy API collect

# 2.4 Sponsors



# 2.5 Community

Polars has a very active community with frequent releases (approximately weekly). Below are some of the top contributors to the project:



# 2.6 Contribute

Thanks for taking the time to contribute! We appreciate all contributions, from reporting bugs to implementing new features. If you're unclear on how to proceed read our contribution guide or contact us on discord.

# 2.7 License

This project is licensed under the terms of the MIT license.

# 3. Getting Started

# 3.1 Introduction

This getting started guide is written for new users of Polars. The goal is to provide a quick overview of the most common functionality. For a more detailed explanation, please go to the User Guide



Due to historical reasons the eager API in Rust is outdated. In the future we would like to redesign it as a small wrapper around the lazy API (as is the design in Python / NodeJS). In the examples we will use the lazy API instead with <code>.lazy()</code> and <code>.collect()</code>. For now you can ignore these two functions. If you want to know more about the lazy and eager API go here.


To enable the Lazy API ensure you have the feature flag lazy configured when installing Polars

```
# Cargo.toml
[dependencies]
polars = { version = "x", features = ["lazy", ...]}
```

Because of the ownership ruling in Rust we can not reuse the same DataFrame multiple times in the examples. For simplicity reasons we call clone() to overcome this issue. Note that this does not duplicate the data but just increments a pointer (Arc).

# 3.2 Installation

Polars is a library and installation is as simple as invoking the package manager of the corresponding programming language.



# 3.2.1 Importing

To use the library import it into your project

```
import polars as pl

use polars::prelude::*;

// esm
import pl from 'nodejs-polars';

// require
const pl = require('nodejs-polars');
```

# 3.3 Series & DataFrames

The core base data structures provided by Polars are  ${\tt Series}$  and  ${\tt DataFrames}$  .

# 3.3.1 Series

Series are a 1-dimensional data structure. Within a series all elements have the same data type (e.g. int, string). The snippet below shows how to create a simple named Series object. In a later section of this getting started guide we will learn how to read data from external sources (e.g. files, database), for now lets keep it simple.







# API Series

```
import polars as pl
s = pl.Series("a", [1, 2, 3, 4, 5])
print(s)
```

#### API Series

```
use chrono::prelude::*;
let s = Series::new("a", [1, 2, 3, 4, 5]);
println!("{}",s);
```

#### API Series

```
const pl = require("nodejs-polars");
var s = pl.Series("a", [1, 2, 3, 4, 5]);
console.log(s);
```

#### Methods

Although it is more common to work directly on a DataFrame object, Series implement a number of base methods which make it easy to perform transformations. Below are some examples of common operations you might want to perform. Note that these are for illustration purposes and only show a small subset of what is available.

#### Aggregations

Series out of the box supports all basic aggregations (e.g. min, max, mean, mode, ...).



```
R Rust
```

```
(js) NodeJS
```

```
API min API max
```

```
s = pl.Series("a", [1, 2, 3, 4, 5])
print(s.min())
print(s.max())
```

API min API max

```
let s = Series::new("a", [1, 2, 3, 4, 5]);
// The use of generics is necessary for the type system
println!("{}",s.min::<u64>().unwrap());
println!("{}",s.max::<u64>().unwrap());
```

API min API max

```
var s = pl.Series("a", [1, 2, 3, 4, 5]);
console.log(s.min());
console.log(s.max());
```

```
1
5
```

#### String

There are a number of methods related to string operations in the StringNamespace. These only work on Series with the Datatype Utf8.







API replace

```
s = pl.Series("a", ["polar", "bear", "arctic", "polar fox", "polar bear"])
s2 = s.str.replace("polar", "pola")
print(s2)
// This operation is not directly available on the Series object yet, only on the DataFrame
```

#### API replace

```
var s = pl.Series("a", ["polar", "bear", "arctic", "polar fox", "polar bear"]);
var s2 = s.str.replace("polar", "pola");
console.log(s2);
```

```
shape: (5,)
Series: 'a' [str]
[
    "pola"
    "bear"
    "arctic"
    "pola fox"
    "pola bear"
]
```

#### Datetime

Similar to strings, there is a seperate namespace for datetime related operations in the DateLikeNameSpace. These only work on Series with DataTypes related to dates.







# API day

```
from datetime import datetime

start = datetime(2001, 1, 1)
stop = datetime(2001, 1, 9)
s = pl.date_range(start, stop, interval="2d", eager=True)
s.dt.day()
print(s)
```

// This operation is not directly available on the Series object yet, only on the DataFrame

# API day

```
var s = pl.Series("a", [
  new Date(2001, 1, 1),
  new Date(2001, 1, 3),
  new Date(2001, 1, 5),
  new Date(2001, 1, 7),
  new Date(2001, 1, 7),
  new Date(2001, 1, 9),
]);
var s2 = s.date.day();
console.log(s2);
```

```
shape: (5,)
Series: 'date' [datetime[µs]]
[
    2001-01-01 00:00:00
    2001-01-05 00:00:00
    2001-01-07 00:00:00
    2001-01-09 00:00:00
]
```

# 3.3.2 DataFrame

A DataFrame is a 2-dimensional data structure that is backed by a Series, and it could be seen as an abstraction of on collection (e.g. list) of Series. Operations that can be executed on DataFrame are very similar to what is done in a SQL like query. You can GROUP BY, JOIN, PIVOT, but also define custom functions. In the next pages we will cover how to perform these transformations.







# API DataFrame

# API DataFrame

#### API DataFrame

```
let df = p1.DataFrame({
  integer: [1, 2, 3, 4, 5],
  date: [
    new Date(2022, 1, 1, 0, 0),
    new Date(2022, 1, 2, 0, 0),
    new Date(2022, 1, 3, 0, 0),
    new Date(2022, 1, 4, 0, 0),
    new Date(2022, 1, 4, 0, 0),
    new Date(2022, 1, 5, 0, 0),
    l,
    float: [4.0, 5.0, 6.0, 7.0, 8.0],
});
console.log(df);
```

# Viewing data

This part focuses on viewing data in a DataFrame . We will use the DataFrame from the previous example as a starting point.

HEAD

The head function shows by default the first 5 rows of a DataFrame. You can specify the number of rows you want to see (e.g. df.head(10)).







API head


print(df.head(3))

API head

println!("{}",df.head(Some(3)));

API head

console.log(df.head(3));



TAIL

The tail function shows the last 5 rows of a DataFrame . You can also specify the number of rows you want to see, similar to head .







API tail


print(df.tail(3))

API tail

println!("{}",df.tail(Some(3)));

API tail

console.log(df.tail(3));



#### SAMPLE

If you want to get an impression of the data of your DataFrame, you can also use sample. With sample you get an n number of random rows from the DataFrame.

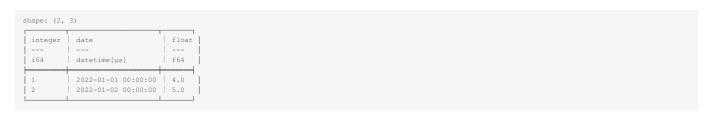


HFI Sampie\_ii

println!("{}",df.sample\_n(2, false, true, None)?);

API sample

console.log(df.sample(2));



# DESCRIBE

 ${\tt Describe} \ \ returns \ summary \ statistics \ of \ your \ \ {\tt DataFrame} \ . \ It \ will \ provide \ several \ quick \ statistics \ if \ possible.$ 



API describe


print(df.describe())

API describe · P. Available on feature describe

println!("{}",df.describe(None));

API describe

console.log(df.describe());



# 3.4 Reading & Writing

Polars supports reading & writing to all common files (e.g. csv, json, parquet), cloud storage (S3, Azure Blob, BigQuery) and databases (e.g. postgres, mysql). In the following examples we will show how to operate on most common file formats. For the following dataframe







#### API DataFrame

#### API DataFrame

#### API DataFrame

```
let df = pl.DataFrame({
  integer: [1, 2, 3],
  date: [
    new Date(2022, 1, 1, 0, 0),
    new Date(2022, 1, 2, 0, 0),
    new Date(2022, 1, 3, 0, 0),
    ],
  float: [4.0, 5.0, 6.0],
]);
console.log(df);
```

```
shape: (3, 3)

integer | date | | float |
--- | --- | | |
i64 | datetime[µs] | f64 |

1 | 2022-01-01 00:00:00 | 4.0 |
2 | 2022-01-02 00:00:00 | 5.0 |
3 | 2022-01-03 00:00:00 | 6.0 |
```

CSV

Polars has its own fast implementation for csv reading with many flexible configuration options.

```
Python
```

```
Rust
```



```
API read_csv API write_csv
```

```
df.write_csv("output.csv")
df_csv = pl.read_csv("output.csv")
print(df_csv)
```

```
API CsvReader API CsvWriter Available on feature csv
```

```
let mut file = File::create("output.csv").expect("could not create file");
CsvWriter::new(&mut file).has_header(true).with_delimiter(b',').finish(&mut df);
let df_csv = CsvReader::from_path("output.csv")?.infer_schema(None).has_header(true).finish()?;
println!("{}",df_csv);
```

# API readCSV API writeCSV

```
df.writeCSV("output.csv");
var df_csv = pl.readCSV("output.csv");
console.log(df_csv);
```

As we can see above, Polars made the datetimes a string. We can tell Polars to parse dates, when reading the csv, to ensure the date becomes a datetime. The example can be found below:







#### API read\_csv

```
df_csv = pl.read_csv("output.csv", try_parse_dates=True)
print(df_csv)
```

API CsvReader · Available on feature csv

```
let mut file = File::create("output.csv").expect("could not create file");
CsvWriter::new(&mut file).has_header(true).with_delimiter(b',').finish(&mut df);
let df_csv = CsvReader::from_path("output.csv")?.infer_schema(None).has_header(true).with_parse_dates(true).finish()?;
println!("{}",df_csv);
```

# API readCSV

```
var df_csv = pl.readCSV("output.csv", { parseDates: true });
console.log(df_csv);
```

```
JSON
                                            (JS) NodeJS
     Python
                         R Rust
  API read_json API write_json
   df.write_json("output.json")
df_json = pl.read_json("output.json")
    print(df_json)
  API JsonReader : API JsonWriter : Available on feature json
    let mut file = File::create("output.json").expect("could not create file");
    JsonWriter::new(&mut file).finish(&mut df);

let mut f = File::open("output.json")?;

let df_json = JsonReader::new(f).with_json_format(JsonFormat::JsonLines).finish()?;
  API readJSON API writeJSON
    df.writeJSON("output.json", { format: "json" });
let df_json = pl.readJSON("output.json");
    console.log(df_json);
    shape: (3, 3)
                                             float
     integer
                  date
     i64
                  \texttt{datetime}[\mu \texttt{s}]
                                             f64
                                             4.0
                   2022-01-02 00:00:00
      3
                                             6.0
PARQUET
     Python
                          Rust
                                            (JS) NodeJS
  API read_parquet API write_parquet
   df.write_parquet("output.parquet")
df_parquet = pl.read_parquet("output.parquet")
print(df_parquet)
  API ParquetReader API ParquetWriter Available on feature parquet
    let mut file = File::create("output.parquet").expect("could not create file");
    ParquetWriter::new(&mut file).finish(&mut df);
```

	ntln!("{}",df_	-	:);
API	readParquet	·API	writeParquet

let mut f = File::open("output.parquet")?;
let df\_parquet = ParquetReader::new(f).finish()?;

df.writeParquet("output.parquet");
let df\_parquet = pl.readParquet("output.parquet");
console.log(df\_parquet);

hape: (3, 3)			
integer	date	float	
i64	datetime[µs]	f64	
1	2022-01-01 00:00:00	4.0	
2	2022-01-02 00:00:00	5.0	
3	2022-01-03 00:00:00	6.0	

To see more examples and other data formats go to the User Guide, section IO.

# 3.5 Expressions

Expressions are the core strength of Polars. The expressions offer a versatile structure that both solves easy queries and is easily extended to complex ones. Below we will cover the basic components that serve as building block (or in Polars terminology contexts) for all your queries:

- select
- filter
- with\_columns
- groupby

To learn more about expressions and the context in which they operate, see the User Guide sections: Contexts and Expressions.

#### Select statement

To select a column we need to do two things. Define the DataFrame we want the data from. And second, select the data that we need. In the example below you see that we select col('\*'). The asterisk stands for all columns.







API select

```
df.select(pl.col("*"))
```

API select

```
let out = df.clone().lazy().select([col("*")]).collect()?;
println!("{}",out);
```

API select

df.select(pl.col("\*"));

```
shape: (8, 4)
        b
                                          d
                                          f64
 i64
        f64
                   datetime[us]
0
        0.577013
                   2022-12-01 00:00:00
                   2022-12-02 00:00:00
        0.114686
                   2022-12-03 00:00:00
        0.612896
                                          NaN
        0.342322
                   2022-12-04 00:00:00
                                          NaN
        0.185987
                   2022-12-05 00:00:00
                                          0.0
                   2022-12-06 00:00:00
        0.286651
                   2022-12-07 00:00:00
                                          -42.0
                   2022-12-08 00:00:00
        0.646312
                                          null
```

You can also specify the specific columns that you want to return. There are two ways to do this. The first option is to create a list of column names, as seen below.







API select

```
df.select(pl.col(["a", "b"]))
```

API select

```
let out = df.clone().lazy().select([col("a"), col("b")]).collect()?;
println!("{}",out);
```

API select

```
df.select(pl.col(["a", "b"]));
```

The second option is to specify each column within a  $\ensuremath{\mathtt{list}}$  in the  $\ensuremath{\mathtt{select}}$  statement. This option is shown below.







API select

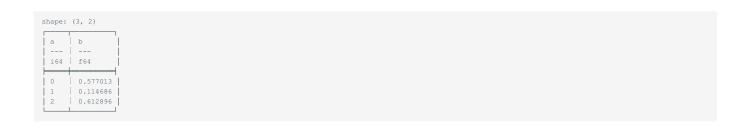
```
df.select([pl.col("a"), pl.col("b")]).limit(3)
```

API select

```
let out = df.clone().lazy().select([col("a"), col("b")]).limit(3).collect()?;
println!("{}",out);
```

API select

```
df.select([pl.col("a"), pl.col("b")]).limit(3);
```



If you want to exclude an entire column from your view, you can simply use exclude in your select statement.







API select

```
df.select([pl.exclude("a")])
```

API select

```
let out = df.clone().lazy().select([col("*").exclude(["a"])]).collect()?;
println!("{}",out);
```

API select

```
df.select([pl.exclude("a")]);
```

```
shape: (8, 3)
 b
                С
                                             d
 f64
                 datetime[us]
                                             f64
 0.577013
                2022-12-01 00:00:00
2022-12-02 00:00:00
                                             1.0
 0.114686
  0.612896
                 2022-12-03 00:00:00
 0.342322
                2022-12-04 00:00:00
2022-12-05 00:00:00
                                             NaN
 0.185987
                                             0.0
  0.376874
                2022-12-06 00:00:00
 0.286651
                2022-12-07 00:00:00
2022-12-08 00:00:00
                                             -42.0
  0.646312
```

# Filter

 $The \ \, \texttt{filter} \ \, \texttt{option} \ \, \texttt{allows} \ \, \texttt{us} \ \, \texttt{to} \ \, \texttt{create} \ \, \texttt{a} \ \, \texttt{subset} \ \, \texttt{of} \ \, \texttt{the} \ \, \texttt{DataFrame} \ \, \texttt{as} \ \, \texttt{earlier} \ \, \texttt{and} \ \, \texttt{we} \ \, \texttt{filter} \ \, \texttt{between} \ \, \texttt{two} \ \, \texttt{specified} \ \, \texttt{dates}.$ 







API filter

```
df.filter(
    pl.col("c").is_between(datetime(2022, 12, 2), datetime(2022, 12, 8)),
)
```

API filter

// TODO

API filter

```
df.filter(pl.col("c").gt(new Date(2022, 12, 2)).lt(new Date(2022, 12, 8)));
```

```
shape: (7, 4)
 a
         b
                       С
                                                d
 i64
         f64
                       \texttt{datetime}[\mu \texttt{s}]
                                                f64
         0.114686
                       2022-12-02 00:00:00
         0.612896
                       2022-12-03 00:00:00
                                                NaN
 3
         0.342322
                       2022-12-04 00:00:00
                                                NaN
         0.185987
                       2022-12-05 00:00:00
                                                0.0
                      2022-12-06 00:00:00
2022-12-07 00:00:00
         0.376874
                                                 -5.0
         0.286651
                                                -42.0
         0.646312
                       2022-12-08 00:00:00
                                                null
```

With filter you can also create more complex filters that include multiple columns.







API filter

```
df.filter((pl.col("a") <= 3) & (pl.col("d").is_not_nan()))</pre>
```

#### API filter

```
let out = df.clone().lazy().filter(col("a").lt_eq(3).and(col("d").is_not_null())).collect()?;
println!("{}",out);
```

#### API filter

```
\texttt{df.filter(pl.col("a").ltEq(3).and(pl.col("d").isNotNull()));}
```

```
shape: (2, 4)

a | b | c | d |
--- | --- | | --- |
i64 | f64 | datetime[µs] | f64 |

0 | 0.577013 | 2022-12-01 00:00:00 | 1.0 |
1 | 0.114686 | 2022-12-02 00:00:00 | 2.0 |
```

#### With columns

with\_columns allows you to create new columns for you analyses. We create two new columns e and b+42. First we sum all values from column b and store the results in column e. After that we add 42 to the values of b. Creating a new column b+42 to store these results.







# API with\_columns

```
df.with columns([pl.col("b").sum().alias("e"), (pl.col("b") + 42).alias("b+42")])
```

# API with columns

# API withColumns

```
df.withColumns([
  pl.col("b").sum().alias("e"),
  pl.col("b").plus(42).alias("b+42"),
]);
```

```
shape: (8, 6)
 а
                    С
                                           d
 i64
        f64
                                           f64
                                                    f64
                                                                f64
                    datetime[us]
 0
        0.577013
                                                    3.142741
                                                                42.577013
        0.114686
                                           2.0
                                                    3.142741
                                                                42.114686
        0.612896
                    2022-12-03 00:00:00
                                                    3.142741
 3
        0.342322
                    2022-12-04 00:00:00
                                           NaN
                                                    3.142741
                                                                42.342322
        0.185987
                    2022-12-05 00:00:00
                                                    3.142741
                                           0.0
                                                                42.185987
        0.376874
                    2022-12-06 00:00:00
                                           -5.0
```

# Groupby

We will create a new DataFrame for the Groupby functionality. This new DataFrame will include several 'groups' that we want to groupby.







#### API DataFrame

# API DataFrame

# API DataFrame

```
df2 = pl.DataFrame({
    x: [...Array(8).keys()],
    y: ["A", "A", "B", "B", "C", "X", "X"],
});
```







# API groupby

```
df2.groupby("y", maintain_order=True).count()
```

# API groupby

```
let out = df2.clone().lazy().groupby(["y"]).agg([count()]).collect()?;
println!("{}",out);
```

# API groupBy

```
df2.groupBy("y").count();
console.log(df2);
```







# API groupby

# API groupby

```
let out = df2.clone().lazy().groupby(["y"]).agg([
    col("*").count().alias("count"),
    col("*").sum().alias("sum"),
]).collect()?;
println!("{}",out);
```

# API groupBy

```
df2
.groupBy("y")
.agg(pl.col("*").count().alias("count"), pl.col("*").sum().alias("sum"));
```

# **Combining operations**

Below are some examples on how to combine operations to create the DataFrame you require.







```
API select API with_columns
```

```
df_x = df.with_columns((pl.col("a") * pl.col("b")).alias("a * b")).select(
     [pl.all().exclude(["c", "d"])]
)
print(df_x)
```

#### API select API with\_columns

```
let out = df.clone().lazy().with_columns({
    (col("a") * col("b")).alias("a * b")
]).select([
    col("*").exclude(["c","d"])
]).collect()?;
println!("{}",out);
```

#### API select API withColumns

```
df_x = df
.withColumns(pl.col("a").mul(pl.col("b")).alias("a * b"))
.select([pl.all().exclude(["c", "d"])]);
console.log(df_x);
```

```
shape: (8, 3)
                     a * b
a
        b
 i64
        f64
                     f64
0
        0.577013
                     0.0
1 2 3
        0.114686
                     0.114686
                   1.225792
        0.612896
        0.342322
        0.185987 | 0.743947
      0.376874 | 1.884371 | 0.286651 | 1.719908 |
```

```
7 | 0.646312 | 4.524185 |
```







API select API with\_columns

```
df_y = df.with_columns([(pl.col("a") * pl.col("b")).alias("a * b")]).select(
        [pl.all().exclude("d")]
)
print(df_y)
```

API select API with\_columns

API select API withColumns

```
df_y = df
.withColumns([pl.col("a").mul(pl.col("b")).alias("a * b")])
.select([pl.all().exclude("d")]);
console.log(df_y);
```

```
shape: (8, 4)
           b
                                                         a * b
                                                         f64
i64
           f64
                           datetime[us]
 0
                        2022-12-01 00:00:00 2022-12-02 00:00:00
           0.577013
                                                          0.0
                                                          0.114686
 1
           0.114686
 2
                           2022-12-03 00:00:00
           0.342322
                        2022-12-04 00:00:00 2022-12-05 00:00:00
                                                        1.026967
                         2022-12-06 00:00:00

    0.286651
    2022-12-07
    00:00:00
    1.719908

    0.646312
    2022-12-08
    00:00:00
    4.524185
```

# 3.6 Combining DataFrames

There are two ways DataFrame s can be combined depending on the use case: join and concat.

# 3.6.1 Join

Polars supports all types of join (e.g. left, right, inner, outer). Let's have a closer look on how to join two DataFrames into a single DataFrame. Our two DataFrames both have an 'id'-like column: a and x. We can use those columns to join the DataFrames in this example.







#### API join

#### API join

# API join

```
df = pl.DataFrame({
    a: [...Array(8).keys()],
    b: Array.from({ length: 8 }, () => Math.random()),
    d: [1, 2.0, null, null, 0, -5, -42, null],
});

df2 = pl.DataFrame({
    x: [...Array(8).keys()],
    y: ["A", "A", "A", "B", "C", "X", "X"],
});
joined = df.join(df2, { leftOn: "a", rightOn: "x" });
console.log(joined);
```

```
shape: (8, 4)
а
                           У ...
                   f64
       f64
                           str
 i64
0
       0.253982
                   1.0
                           Α
        0.396669
        0.131986
       0.143672
                   NaN
       0.356455
                   0.0
        0.111859
       0.399095
                  null
                          Х
       0.150452
```

To see more examples with other types of joins, go the User Guide.

# 3.6.2 Concat

We can also concatenate two DataFrames. Vertical concatenation will make the DataFrame longer. Horizontal concatenation will make the DataFrame wider. Below you can see the result of an horizontal concatenation of our two DataFrames.







#### API hstack

```
stacked = df.hstack(df2)
print(stacked)
```

# API hstack

```
let stacked = df.hstack(df2.get_columns())?;
println!("{}",stacked);
```

#### API hstack

stacked = df.hstack(df2);
console.log(stacked);

#### shape: (8, 5) d i64 f64 i64 f64 0 0.253982 1.0 0.396669 0.131986 0.143672 NaN В NaN B 0.356455 -5.0 -42.0 0.111859 0.399095 0.150452

# 4. User Guide

# 4.1 Introduction

This User Guide is an introduction to the Polars DataFrame library. Its goal is to introduce you to Polars by going through examples and comparing it to other solutions. Some design choices are introduced here. The guide will also introduce you to optimal usage of Polars.

Even though Polars is completely written in Rust (no runtime overhead!) and uses Arrow -- the native arrow2 Rust implementation -- as its foundation, the examples presented in this guide will be mostly using its higher-level language bindings. Higher-level bindings only serve as a thin wrapper for functionality implemented in the core library.

For Pandas users, our Python package will offer the easiest way to get started with Polars.

# **Philosophy**

The goal of Polars is to provide a lightning fast DataFrame library that:

- Utilizes all available cores on your machine.
- Optimizes queries to reduce unneeded work/memory allocations.
- Handles datasets much larger than your available RAM.
- Has an API that is consistent and predictable.
- Has a strict schema (data-types should be known before running the query).

Polars is written in Rust which gives it C/C++ performance and allows it to fully control performance critical parts in a query engine.

As such Polars goes to great lengths to:

- Reduce redundant copies.
- Traverse memory cache efficiently.
- Minimize contention in parallelism.
- Process data in chunks.
- Reuse memory allocations.

# 4.2 Installation

Polars is a library and installation is as simple as invoking the package manager of the corresponding programming language.

# 4.2.1 Importing

To use the library import it into your project

```
import polars as pl
use polars::prelude::*;
// esm
import pl from 'nodejs-polars';
// require
const pl = require('nodejs-polars');
```

# 4.2.2 Feature Flags

By using the above command you install the core of Polars onto your system. However depending on your use case you might want to install the optional dependencies as well. These are made optional to minimize the footprint. The flags are different depending on the programming language. Throughout the user guide we will mention when a functionality is used that requires an additional dependency.

# Python

# For example pip install polars[numpy, fsspec]		
Tag	Description	
all	Install all optional dependencies (all of the following)	
pandas	Install with Pandas for converting data to and from Pandas Dataframes/Series	
numpy	Install with numpy for converting data to and from numpy arrays	
pyarrow	Reading data formats using PyArrow	
fsspec	Support for reading from remote file systems	
connectorx	Support for reading from SQL databases	
xlsx2csv	Support for reading from Excel files	
deltalake	Support for reading from Delta Lake Tables	
timezone	Timezone support, only needed if 1. you are on Python < 3.9 and/or 2. you are on Windows, otherwise no dependencies will be installed	

# Rust

```
# Cargo.toml
[dependencies]
polars = { version = "0.26.1", features = ["lazy", "temporal", "describe", "json", "parquet", "dtype-datetime"]}
```

The opt-in features are:

- · Additional data types:
- · dtype-date
- dtype-datetime
- · dtype-time
- dtype-duration
- dtype-i8
- dtype-i16
- dtype-u8
- dtype-u16
- dtype-categorical
- dtype-struct
- performant Longer compile times more fast paths.
- lazy Lazy API
- lazy regex Use regexes in column selection
- dot\_diagram Create dot diagrams from lazy logical plans.
- sql Pass SQL queries to polars.
- streaming Be able to process datasets that are larger than RAM.
- random Generate arrays with randomly sampled values
- ndarray Convert from DataFrame to ndarray
- temporal Conversions between Chrono and Polars for temporal data types
- timezones Activate timezone support.
- strings Extra string utilities for Utf8Chunked
- string justify zfill, ljust, rjust
- string\_from\_radix parse\_int
- object Support for generic ChunkedArrays called ObjectChunked<T> (generic over T). These are downcastable from Series through the Any trait.
- Performance related:
- $\bullet$  nightly Several nightly only features such as SIMD and specialization.
- performant more fast paths, slower compile times.
- bigidx Activate this feature if you expect >> 2^32 rows. This has not been needed by anyone. This allows polars to scale up way beyond that by using u64 as an index. Polars will be a bit slower with this feature activated as many data structures are less cache efficient.
- cse Activate common subplan elimination optimization
- IO related:
- serde Support for serde serialization and descrialization. Can be used for JSON and more serde supported serialization formats.
- serde-lazy Support for serde serialization and descrialization. Can be used for JSON and more serde supported serialization formats.
- parquet Read Apache Parquet format
- | json JSON serialization
- ipc Arrow's IPC format serialization
- decompress Automatically infer compression of csvs and decompress them. Supported compressions: zip gzip
- DataFrame operations:
- dynamic\_groupby Groupby based on a time window instead of predefined keys. Also activates rolling window group by operations.
- $\bullet$  sort\_multiple Allow sorting a DataFrame on multiple columns
- rows Create DataFrame from rows and extract rows from DataFrames . And activates pivot and transpose operations
- join\_asof Join ASOF, to join on nearest keys instead of exact equality match.
- $\bullet$   ${\tt cross\_join}$  Create the cartesian product of two DataFrames.
- semi anti join SEMI and ANTI joins.
- groupby\_list Allow groupby operation on keys of type List.
- row\_hash Utility to hash DataFrame rows to UInt64Chunked
- diagonal\_concat Concat diagonally thereby combining different schemas.
- $\verb| horizontal_concat| Concat horizontally and extend with null values if lengths don't match$
- dataframe arithmetic Arithmetic on (Dataframe and DataFrames) and (DataFrame on Series)
- partition\_by Split into multiple DataFrames partitioned by groups.

- Series / Expression operations:
- is in Check for membership in Series
- zip with Zip two Series/ ChunkedArrays
- round series round underlying float types of Series.
- repeat\_by [Repeat element in an Array N times, where N is given by another array.
- is first Check if element is first unique value.
- is\_last Check if element is last unique value.
- checked\_arithmetic checked arithmetic/ returning None on invalid operations.
- dot product Dot/inner product on Series and Expressions.
- concat str Concat string data in linear time.
- reinterpret Utility to reinterpret bits to signed/unsigned
- take\_opt\_iter Take from a Series with Iterator<Item=Option<usize>>
- mode Return the most occurring value(s)
- cum\_agg cumsum, cummin, cummax aggregation.
- rolling\_window rolling window functions, like rolling\_mean
- interpolate interpolate None values
- extract\_jsonpath Run jsonpath queries on Utf8Chunked
- list List utils.
- · list\_take take sublist by multiple indices
- rank Ranking algorithms.
- moment kurtosis and skew statistics
- ewma Exponential moving average windows
- · abs Get absolute values of Series
- arange Range operation on Series
- product Compute the product of a Series.
- diff diff operation.
- pct change Compute change percentages.
- unique\_counts Count unique values in expressions.
- log Logarithms for Series.
- list\_to\_struct Convert List to Struct dtypes.
- list count Count elements in lists.
- list\_eval Apply expressions over list elements.
- cumulative eval Apply expressions over cumulatively increasing windows.
- arg where Get indices where condition holds.
- search sorted Find indices where elements should be inserted to maintain order.
- date\_offset Add an offset to dates that take months and leap years into account.
- trigonometry Trigonometric functions.
- sign Compute the element-wise sign of a Series.
- propagate\_nans NaN propagating min/max aggregations.
- DataFrame pretty printing
- fmt Activate DataFrame formatting

# 4.3 Concepts

# 4.3.1 Data types

Polars is entirely based on Arrow data types and backed by Arrow memory arrays. This makes data processing cache-efficient and well-supported for Inter Process Communication. Most data types follow the exact implementation from Arrow, with the exception of Utf8 (this is actually LargeUtf8), Categorical, and Object (support is limited). The data types are:

Group	Туре	Details
Numeric	Int8	8-bit signed integer.
	Int16	16-bit signed integer.
	Int32	32-bit signed integer.
	Int64	64-bit signed integer.
	UInt8	8-bit unsigned integer.
	UInt16	16-bit unsigned integer.
	UInt32	32-bit unsigned integer.
	UInt64	64-bit unsigned integer.
	Float32	32-bit floating point.
	Float64	64-bit floating point.
Nested	Struct	A struct array is represented as a <code>Vec<series></series></code> and is useful to pack multiple/heterogenous values in a single column.
	List	A list array contains a child array containing the list values and an offset array. (this is actually Arrow LargeList internally).
Temporal	Date	Date representation, internally represented as days since UNIX epoch encoded by a 32-bit signed integer.
	Datetime	Datetime representation, internally represented as microseconds since UNIX epoch encoded by a 64-bit signed integer.
	Duration	A timedelta type, internally represented as microseconds. Created when subtracting <code>Date/Datetime</code> .
	Time	Time representation, internally represented as nanoseconds since midnight.
Other	Boolean	Boolean type effectively bit packed.
	Utf8	String data (this is actually Arrow LargeUtf8 internally).
	Binary	Store data as bytes.
	Object	A limited supported data type that can be any value.
	Categorical	A categorical encoding of a set of strings.

To learn more about the internal representation of these data types, check the  ${\tt Arrow}$  columnar format.

# 4.3.2 Data Structures

The core base data structures provided by Polars are  ${\tt Series}$  and  ${\tt DataFrames}$  .

# Series

Series are a 1-dimensional data structure. Within a series all elements have the same Data Type . The snippet below shows how to create a simple named Series object.







# API Series

```
import polars as pl
s = pl.Series("a", [1, 2, 3, 4, 5])
print(s)
```

#### API Series

```
use chrono::prelude::*;
let s = Series::new("a", [1, 2, 3, 4, 5]);
println!("{}",s);
```

# API Series

```
const pl = require("nodejs-polars");
var s = pl.Series("a", [1, 2, 3, 4, 5]);
console.log(s);
```

#### **DataFrame**

A DataFrame is a 2-dimensional data structure that is backed by a Series, and it can be seen as an abstraction of a collection (e.g. list) of Series. Operations that can be executed on a DataFrame are very similar to what is done in a SQL like query. You can GROUP BY, JOIN, PIVOT, but also define custom functions.







### API DataFrame

### API DataFrame

### API DataFrame

```
let df = pl.DataFrame({
   integer: [1, 2, 3, 4, 5],
   date: [
   new Date(2022, 1, 1, 0, 0),
   new Date(2022, 1, 2, 0, 0),
   new Date(2022, 1, 3, 0, 0),
   new Date(2022, 1, 4, 0, 0),
   new Date(2022, 1, 4, 0, 0),
   new Date(2022, 1, 5, 0, 0),
   ],
   float: [4.0, 5.0, 6.0, 7.0, 8.0],
});
console.log(df);
```

```
shape: (5, 3)

integer | date | | float |
--- | --- |
i64 | datetime[µs] | f64

1 | 2022-01-01 00:00:00 | 4.0 |
2 | 2022-01-02 00:00:00 | 5.0 |
3 | 2022-01-03 00:00:00 | 6.0 |
4 | 2022-01-04 00:00:00 | 7.0 |
5 | 2022-01-05 00:00:00 | 8.0
```

### VIEWING DATA

This part focuses on viewing data in a DataFrame . We will use the DataFrame from the previous example as a starting point.

Head

The head function shows by default the first 5 rows of a DataFrame. You can specify the number of rows you want to see (e.g. df.head(10)).







API head

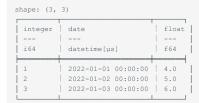
print(df.head(3))

API head

println!("{}",df.head(Some(3)));

API head

console.log(df.head(3));



Tail

The tail function shows the last 5 rows of a DataFrame. You can also specify the number of rows you want to see, similar to head.







API tail


print(df.tail(3))

API tail

println!("{}",df.tail(Some(3)));

API tail

console.log(df.tail(3));



### Sample

If you want to get an impression of the data of your DataFrame, you can also use sample. With sample you get an n number of random rows from the DataFrame.







API sample

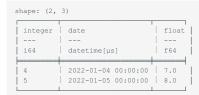
print(df.sample(2))

API sample n

println!("{}",df.sample\_n(2, false, true, None)?);

API sample

console.log(df.sample(2));



### Describe

 ${\tt Describe} \ \ returns \ summary \ statistics \ of \ your \ \ {\tt DataFrame} \ . \ It \ will \ provide \ several \ quick \ statistics \ if \ possible.$ 







API describe


print(df.describe())

API describe · Available on feature describe

println!("{}",df.describe(None));

API describe

console.log(df.describe());



### 4.3.3 Contexts

Polars has developed its own Domain Specific Language (DSL) for transforming data. The language is very easy to use and allows for complex queries that remain human readable. The two core components of the language are Contexts and Expressions, the latter we will cover in the next section.

A context, as implied by the name, refers to the context in which an expression needs to be evaluated. There are three main contexts 1:

- 1. Selection: df.select([..]) , df.with\_columns([..])
- 2. Filtering: df.filter()
- 3. Groupby / Aggregation: df.groupby(..).agg([..])

The examples below are performed on the following  ${\tt DataFrame}$  :







### API DataFrame

### API DataFrame

```
use rand::{thread_rng, Rng};
let mut arr = [0f64; 5];
thread_rng().fill(&mut arr);
let df = df! (
    "nrs" => &[Some(1), Some(2), Some(3), None, Some(5)],
    "names" => &[Some("foo"), Some("ham"), Some("spam"), Some("eggs"), None],
    "random" => &arr,
    "groups" => &["A", "A", "B", "C", "B"],
)?;
println!("{}", &df);
```

# API DataFrame

```
const arr = Array.from({ length: 5 }).map((_) =>
    chance.floating({ min: 0, max: 1 }),
);

let df = pl.DataFrame({
    nrs: [1, 2, 3, null, 5],
    names: ["foo", "ham", "spam", "egg", null],
    random: arr,
    groups: ["A", "A", "B", "C", "B"],
});
console.log(df);
```

```
shape: (5, 4)
 i64
         str
                             str
                  0.154163
 1
                             А
         foo
                  0.74005
         ham
                  0.263315
         spam
         egg
null
 null
                  0.533739
                           В
                  0.014575
```

#### Select

In the select context the selection applies expressions over columns. The expressions in this context must produce Series that are all the same length or have a length of 1.

A Series of a length of 1 will be broadcasted to match the height of the DataFrame. Note that a select may produce new columns that are aggregations, combinations of expressions, or literals.







### API select

```
out = df.select(
  pl.sum("nrs"),
  pl.col("names").sort(),
  pl.col("names").first().alias("first name"),
      (pl.mean("nrs") * 10).alias("10xnrs"),
)
print(out)
```

### API select

### API select

```
let out = df.select(
  pl.col("nrs").sum(),
  pl.col("names").sort(),
  pl.col("names").first().alias("first name"),
  pl.mean("nrs").multiplyBy(10).alias("10xnrs"),
);
console.log(out);
```

```
shape: (5, 4)
 nrs
        names
                  first name
                                10xnrs
 i64
         str
                  str
                                f64
                                27.5
                  foo
                                27.5
 11
         foo
                  foo
                                27.5
 11
         ham
                  foo
 11
                  foo
                                27.5
```

As you can see from the query the select context is very powerful and allows you to perform arbitrary expressions independent (and in parallel) of each other.

Similarly to the select statement there is the with\_columns statement which also is an entrance to the selection context. The main difference is that with\_columns retains the original columns and adds new ones while select drops the original columns.







API with\_columns

```
df = df.with_columns(
    pl.sum("nrs").alias("nrs_sum"),
    pl.col("random").count().alias("count"),
)
print(df)
```

### API with\_columns

```
let out = df
    .clone()
    .lazy()
    .with_columns([
        sum("nrs").alias("nrs_sum"),
        col("random").count().alias("count"),
    ])
    .collect()?;
println!("{}", out);
```

### API withColumns

```
df = df.withColumns(
   pl.col("nrs").sum().alias("nrs_sum"),
   pl.col("random").count().alias("count"),
);
console.log(df);
```

```
shape: (5, 6)
         names
                  random
                              groups
                                        nrs_sum
 i64
          str
                  f64
                              str
                                        i64
                                                   u32
 1
         foo
                  0.154163
         ham
         spam
                  0.263315
                              В
                  0.533739
 null
         egg
                  0.014575
```

# Filter

In the filter context you filter the existing dataframe based on arbritary expression which evaluates to the Boolean data type.







### API filter

```
out = df.filter(pl.col("nrs") > 2)
print(out)
```

### API filter

```
let out = df.clone().lazy().filter(col("nrs").gt(lit(2))).collect()?;
println!("{}", out);
```

# API filter

```
out = df.filter(pl.col("nrs").gt(2));
console.log(out);
```

```
shape: (2, 6)
```

### **Groupby / Aggregation**

In the groupby context expressions work on groups and thus may yield results of any length (a group may have many members).







### API groupby

```
out = df.groupby("groups").agg(
   pl.sum("nrs"),   # sum nrs by groups
   pl.col("random").count().alias("count"),   # count group members
    # sum random where name != null
   pl.col("random").filter(pl.col("names").is_not_null()).sum().suffix("_sum"),
   pl.col("names").reverse().alias(("reversed names")),
)
print(out)
```

### API groupby

### API groupBy

```
out = df.groupBy("groups").agg(
pl
    .col("nrs")
    .sum(), // sum nrs by groups
pl
    .col("random")
    .count()
    .alias("count"), // count group members
// sum random where name != null
pl
    .col("random")
    .filter(pl.col("names").isNotNull())
    .sum()
    .suffix("_sum"),
pl.col("names").reverse().alias("reversed names"),
);
console.log(out);
```

```
shape: (3, 5)
                     count
                               random_sum
                                              reversed names
 groups
            nrs
                               f64
 str
            i64
                     u32
                                              ["ham", "foo"]
[null, "spam"]
 Α
            3
                     2
                               0.894213
            null
                                               ["egg"]
```

As you can see from the result all expressions are applied to the group defined by the <code>groupby</code> context. Besides the standard <code>groupby\_dynamic</code>, and <code>groupby\_rolling</code> are also entrances to the groupby context.

1. There are additional List and SQL contexts which are covered later in this guide. But for simplicity, we leave them out of scope for now.  $\leftarrow$ 

### 4.3.4 Expressions

Polars has a powerful concept called expressions that is central to its very fast performance.

Expressions are at the core of many data science operations:

- taking a sample of rows from a column
- multiplying values in a column
- extracting a column of years from dates
- · convert a column of strings to lowercase
- and so on!

However, expressions are also used within other operations:

- taking the mean of a group in a groupby operation
- calculating the size of groups in a groupby operation
- taking the sum horizontally across columns

Polars performs these core data transformations very quickly by:

- automatic query optimization on each expression
- automatic parallelization of expressions on many columns

Polars expressions are a mapping from a series to a series (or mathematically Fn (Series) -> Series). As expressions have a Series as an input and a Series as an output then it is straightforward to do a sequence of expressions (similar to method chaining in Pandas).

### **Examples**

The following is an expression:



The snippet above says:

- 1. Select column "foo"
- 2. Then sort the column (not in reversed order)
- 3. Then take the first two values of the sorted output

The power of expressions is that every expression produces a new expression, and that they can be *piped* together. You can run an expression by passing them to one of Polars execution contexts.

Here we run two expressions by running  ${\, {\tt df.select}} :$ 







API select

```
df.select(pl.col("foo").sort().head(2), pl.col("bar").filter(pl.col("foo") == 1).sum())
```

API select

```
df.clone().lazy().select([
   col("foo").sort(Default::default()).head(Some(2)),
   col("bar").filter(col("foo").eq(lit(1))).sum(),
]).collect()?;
```

API select

```
df.select(
  pl.col("foo").sort().head(2),
  pl.col("bar").filter(pl.col("foo").eq(1)).sum(),
);
```

All expressions are run in parallel, meaning that separate Polars expressions are **embarrassingly parallel**. Note that within an expression there may be more parallelization going on.

### Conclusion

This is the tip of the iceberg in terms of possible expressions. There are a ton more, and they can be combined in a variety of ways. This page is intended to get you familiar with the concept of expressions, in the section on expressions we will dive deeper.

# 4.3.5 Lazy / Eager API

Polars supports two modes of operation: lazy and eager. In the eager API the query is executed immediately while in the lazy API the query is only evaluated once it is 'needed'. Deferring the execution to the last minute can have significant performance advantages that is why the Lazy API is preferred in most cases. Let us demonstrate this with an example:







API read\_csv

```
df = pl.read_csv("docs/src/data/iris.csv")
df_small = df.filter(pl.col("sepal_length") > 5)
df_agg = df_small.groupby("species").agg(pl.col("sepal_width").mean())
print(df_agg)
```

API CsvReader ·



```
let df = CsvReader::from_path("docs/src/data/iris.csv").unwrap().finish().unwrap();
let mask = df.column("sepal_width")?.f64()?.gt(5.0);
let df_small = df.filter(&mask)?;
let df_agg = df_small.groupby(["species"])?.select(["sepal_width"]).mean()?;
println!("{}", df_agg);
```

API readCSV

```
df = pl.readCSV("docs/src/data/iris.csv");
df_small = df.filter(pl.col("sepal_length").gt(5));
df_agg = df_small.groupBy("species").agg(pl.col("sepal_width").mean());
console.log(df_agg);
```

In this example we use the eager API to:

- 1. Read the iris dataset.
- 2. Filter the dataset based on sepal length
- 3. Calculate the mean of the sepal width per species

Every step is executed immediately returning the intermediate results. This can be very wasteful as we might do work or load extra data that is not being used. If we instead used the lazy API and waited on execution until all the steps are defined then the query planner could perform various optimizations. In this case:

- Predicate pushdown: Apply filters as early as possible while reading the dataset, thus only reading rows with sepal length greater than 5.
- Projection pushdown: Select only the columns that are needed while reading the dataset, thus removing the need to load additional columns (e.g. petal length & petal width)







### API scan\_csv

```
q = (
    pl.scan_csv("docs/src/data/iris.csv")
    .filter(pl.col("sepal_length") > 5)
    .groupby("species")
    .agg(pl.col("sepal_width").mean())
)

df = q.collect()
```

API LazyCsvReader · Available on feature csv

```
let q = LazyCsvReader::new("docs/src/data/iris.csv")
    .has_header(true)
    .finish()?
    .filter(col("sepal_length").gt(lit(5)))
    .groupby(vec![col("species")])
    .agg([col("sepal_width").mean()]);
let df = q.collect()?;
println!("{}", df);
```

### API scanCSV

```
q = pl
.scanCSV("docs/src/data/iris.csv")
.filter(pl.col("sepal_length").gt(5))
.groupBy("species")
.agg(pl.col("sepal_width").mean());

df = q.collect();
```

These will significantly lower the load on memory & CPU thus allowing you to fit bigger datasets in memory and process faster. Once the query is defined you call collect to inform Polars that you want to execute it. In the section on Lazy API we will go into more details on its implementation.



In many cases the eager API is actually calling the lazy API under the hood and immediately collecting the result. This has the benefit that within the query itself optimization(s) made by the query planner can still take place.

### WHEN TO USE WHICH

In general the lazy API should be preferred unless you are either interested in the intermediate results or are doing exploratory work and don't know yet what your query is going to look like.

### 4.3.6 Streaming API

One additional benefit of the lazy API is that it allows queries to be executed in a streaming manner. Instead of processing the data all-at-once Polars can execute the query in batches allowing you to process datasets that are larger-than-memory.

 $To \ tell \ Polars \ we \ want \ to \ execute \ a \ query \ in \ streaming \ mode \ we \ pass \ the \ \ streaming = True \ argument \ to \ \ collect$ 





API collect

```
q = (
    pl.scan_csv("docs/src/data/iris.csv")
    .filter(pl.col("sepal_length") > 5)
    .groupby("species")
    .agg(pl.col("sepal_width").mean())
)

df = q.collect(streaming=True)
```

API collect · Available on feature streaming

### When is streaming available?

Streaming is still in development. We can ask Polars to execute any lazy query in streaming mode. However, not all lazy operations support streaming. If there is an operation for which streaming is not supported Polars will run the query in non-streaming mode.

Streaming is supported for many operations including:

- filter, slice, head, tail
- with\_columns, select
- groupby
- join
- sort
- explode, melt
- scan\_csv , scan\_parquet , scan\_ipc

# 4.4 Expressions

### 4.4.1 Basic Operators

This section describes how to use basic operators (e.g. addition, substraction) in conjunction with Expressions. We will provide various examples using different themes in the context of the following dataframe.



In Rust and Python it is possible to use the operators directly (as in + - \* / < >) as the language allows operator overloading. For instance, the operator + translates to the .add() method. In NodeJS this is not possible and you must use the methods themselves, in python and rust you can choose which one you prefer.



### API DataFrame

```
shape: (5, 4)
 nrs
                  random
                              groups
 i64
                  f64
                              str
         str
         foo
                  0.154163
         ham
                  0.74005
                  0.263315
         spam
                  0.533739
         null
                  0.014575
                              В
```

### NUMERICAL



## API operators

```
df_numerical = df.select(
    (pl.col("nrs") + 5).alias("nrs + 5"),
    (pl.col("nrs") - 5).alias("nrs - 5"),
    (pl.col("nrs") * pl.col("random")).alias("nrs * random"),
    (pl.col("nrs") / pl.col("random")).alias("nrs / random"),
)
print(df_numerical)
```

### LOGICAL



# API operators

```
df_logical = df.select(
    (pl.col("nrs") > 1).alias("nrs > 1"),
    (pl.col("random") <= 0.5).alias("random < .5"),
    (pl.col("nrs") != 1).alias("nrs != 1"),
    (pl.col("nrs") == 1).alias("nrs == 1"),
    ((pl.col("random") <= 0.5) & (pl.col("nrs") > 1)).alias("and_expr"), # and
    ((pl.col("random") <= 0.5) | (pl.col("nrs") > 1)).alias("or_expr"), # or
)
print(df_logical)
```

nrs > 1	random < .5	nrs != 1	nrs == 1	and_expr	or_expr
bool	bool	bool	bool	bool	bool
false	true	false	true	false	true
true	false	true	false	false	true
true	true	true	false	true	true
null	false	null	null	false	null
	true	true	false	true	true

### 4.4.2 Column Selections

Let's create a dataset to use in this section:



### API DataFrame

```
shape: (3, 7)
  rn
                                                       has_people
                                                                      logged_at
                                         f64
 u32
         i64
                                                       bool
                                                                     datetime[µs]
                str
                           date
                Mars
                                         33.4
                                                       false
                                                                      2022-12-01 00:00:00
                                         2142134
                Earth
                                                       true
                Saturn
```

### **Expression Expansion**

As we've seen in the previous section, we can select specific columns using the pl.col method. It can also select multiple columns - both as a means of convenience, and to *expand* the expression.

This kind of convenience feature isn't just decorative or syntactic sugar. It allows for a very powerful application of DRY principles in your code: a single expression that specifies multiple columns expands into a list of expressions (depending on the DataFrame schema), resulting in being able to select multiple columns + run computation on them!

SELECT ALL, OR ALL BUT SOME

We can select all columns in the DataFrame object by providing the argument \*:



### API all

```
out = df.select(pl.col("*"))

# Is equivalent to
out = df.select(pl.all())
print(out)
```

```
shape: (3, 7)
 rn
         id
               place
                          date
                                        sales
                                                     has_people
                                                                    logged_at
  u32
         i64
                str
                          date
                                        f64
                                                                    datetime[µs]
                                                                    2022-12-01 00:00:00
               Mars
                          2022-01-01
                                        33.4
                                                      false
               Earth
                          2022-01-02
                                        2142134.1
                                                                    2022-12-01 00:00:01
        2
               Saturn
                                        44.7
                                                     false
                                                                    2022-12-01 00:00:02
```

Often, we don't just want to include all columns, but include all while excluding a few. This can be done easily as well:



### API exclude

```
out = df.select(pl.col("*").exclude("logged_at", "rn"))
print(out)
```

id   p	lace	date	sales	has people
i64   s			f64	bool
9 M	lars	2022-01-01	33.4	false
4 E	arth	2022-01-02	2142134.1	true
2   S		2022-01-03	1 44 5	false

### BY MULTIPLE STRINGS

Specifying multiple strings allows expressions to expand to all matching columns:



### API dt.to\_string

```
out = df.select(pl.col("date", "logged_at").dt.to_string("%Y-%h-%d"))
print(out)
```

```
date | logged_at | --- | --- | str | str | 2022-Jan-01 | 2022-Dec-01 | 2022-Jan-03 | 2022-Dec-01 |
```

### BY REGULAR EXPRESSIONS

Multiple column selection is possible by regular expressions also, by making sure to wrap the regex by ^ and \$ to let pl.col know that a regex selection is expected:

# Python

```
out = df.select(pl.col("^.*(as|sa).*$"))
print(out)
```

sales has_peop
F64 bool
33.4 false
2142134.1   true
14.7 false

### BY DATA TYPE

pl.col can select multiple columns using Polars data types:



### API n\_unique

```
out = df.select(pl.col(pl.Int64, pl.UInt32, pl.Boolean).n_unique())
print(out)
```

### Using selectors

Polars also allows for the use of intuitive selections for columns based on their name, dtype or other properties; and this is built on top of existing functionality outlined in col used above. It is recommended to use them by importing and aliasing polars.selectors as cs.

#### BY DTYPE

To select just the integer and string columns, we can do:



### API selectors

```
import polars.selectors as cs

out = df.select(cs.integer(), cs.string())
print(out)
```

### APPLYING SET OPERATIONS

These selectors also allow for set based selection operations. For instance, to select the numeric columns except the first column that indicates row numbers:



### API cs.first API cs.numeric

```
out = df.select(cs.numeric() - cs.first())
print(out)
```

```
shape: (3, 2)

id | sales | --- | --- |
i64 | f64 |
| 9 | 33.4 |
4 | 2142134.1 |
2 | 44.7 |
```

We can also select the row number by name  $\boldsymbol{and}$  any  $\boldsymbol{non}\text{-}numeric$  columns:



API cs.by\_name API cs.numeric

```
out = df.select(cs.by_name("rn") | ~cs.numeric())
print(out)
```

place	date	has_people	logged_at
str	date	bool	datetime[µs]
Mars	2022-01-01	false	2022-12-01 00:00:00
Earth	2022-01-02	true	2022-12-01 00:00:01
Saturn	2022-01-03	false	2022-12-01 00:00:02

### BY PATTERNS AND SUBSTRINGS

Selectors can also be matched by substring and regex patterns:



API cs.contains API cs.matches

```
out = df.select(cs.contains("rn"), cs.matches(".*_.*"))
print(out)
```

### CONVERTING TO EXPRESSIONS

What if we want to apply a specific operation on the selected columns (i.e. get back to representing them as **expressions** to operate upon)? We can simply convert them using as\_expr and then proceed as normal:



API cs.temporal

```
out = df.select(cs.temporal().as_expr().dt.to_string("%Y-%h-%d"))
print(out)
```

```
date | logged_at | --- | str | str | str | 2022-Jan-01 | 2022-Dec-01 | 2022-Jan-03 | 2022-Dec-01 |
```

### DEBUGGING SELECTORS

Polars also provides two helpful utility functions to aid with using selectors: is\_selector and selector\_column\_names:



### API is\_selector

```
from polars.selectors import is_selector

out = cs.temporal()
print(is_selector(out))
```

True

To predetermine the column names that are selected, which is especially useful for a LazyFrame object:



# API selector\_column\_names

```
from polars.selectors import selector_column_names

out = cs.temporal().as_expr().dt.to_string("%Y-%h-%d")
print(selector_column_names(df, out))
```

```
('date', 'logged_at')
```

### 4.4.3 Functions

Polars expressions have a large number of built in functions. These allow you to create complex queries without the need for user defined functions. There are too many to go through here, but we will cover some of the more popular use cases. If you want to view all the functions go to the API Reference for your programming language.

In the examples below we will use the following DataFrame:



### API DataFrame

```
shape: (5, 4)
 nrs
                  random
                              groups
         names
 i64
         str
                  f64
                              str
 1
         foo
                  0.154163
         ham
                  0.74005
                  0.263315
         spam
 null
                  0.533739
         egg
         spam
                  0.014575
                           В
```

### **Column Naming**

By default if you perform an expression it will keep the same name as the original column. In the example below we perform an expression on the nrs column. Note that the output DataFrame still has the same name.

```
Python

df_samename = df.select(pl.col("nrs") + 5)
print(df_samename)

df_samename = df.select(pl.col("nrs") + 5)
print(df_samename)
```

```
nrs | --- | i64 | 6 | 7 | 8 | null | 10 |
```

This might get problematic in the case you use the same column multiple times in your expression as the output columns will get duplicated. For example, the following query will fail.

```
try:
    df_samename2 = df.select(pl.col("nrs") + 5, pl.col("nrs") - 5)
    print(df_samename2)
except Exception as e:
    print(e)
```

```
column with name 'nrs' has more than one occurrences
```

You can change the output name of an expression by using the alias function

```
Python
```

### API alias

In case of multiple columns for example when using all() or col(\*) you can apply a mapping function  $map\_alias$  to change the original column name into something else. In case you want to add a suffix (suffix()) or prefix() these are also built in.



```
API prefix API suffix API map_alias
```

### **Count Unique Values**

There are two ways to count unique values in Polars: an exact methodology and an approximation. The approximation uses the HyperLogLog++ algorithm to approximate the cardinality and is especially useful for very large datasets where an approximation is good enough.

```
Python
```

API n\_unique .API approx\_unique

```
df_alias = df.select(
    pl.col("names").n_unique().alias("unique"),
    pl.approx_unique("names").alias("unique_approx"),
)
print(df_alias)
```

```
shape: (1, 2)

unique | unique_approx |
--- | --- |
u32 | u32 |
4 | 4 |
```

### Conditionals

Polars supports if-else like conditions in expressions with the when, then, otherwise syntax. The predicate is placed in the when clause and when this evaluates to true the then expression is applied otherwise the otherwise expression is applied (row-wise).



### API when

```
df_conditional = df.select(
    pl.col("nrs"),
    pl.when(pl.col("nrs") > 2)
    .then(pl.lit(True))
    .otherwise(pl.lit(False))
    .alias("conditional"),
)
print(df_conditional)
```

nrs	shape: (5,	2)			
	nrs c	conditional	¬ .		
1			İ		
2	164 b	0001			
2	1 f.	false	<b>∃</b> 		
null   false					
	3   t	rue			
5 true	null f	false			
	5   t	rue			

### 4.4.4 Casting

Casting converts the underlying <code>DataType</code> of a column to a new one. Polars uses Arrow to manage the data in memory and relies on the compute kernels in the rust implementation to do the conversion. Casting is available with the <code>cast()</code> method.

The cast method includes a strict parameter that determines how Polars behaves when it encounters a value that can't be converted from the source DataType to the target DataType. By default, strict=True, which means that Polars will throw an error to notify the user of the failed conversion and provide details on the values that couldn't be cast. On the other hand, if strict=False, any values that can't be converted to the target DataType will be quietly converted to null.

### **Numerics**

Let's take a look at the following DataFrame which contains both integers and floating point numbers.



#### API DataFrame

```
shape: (5, 4)
              big_integers
                              floats
                                        floats_with_decimal
 integers
 i64
                              f64
              i64
                                        f64
                              4.0
                                         4.532
                              6.0
                                         6.5
              10000004
                              8.0
                                        8.5
```

To perform casting operations between floats and integers, or vice versa, we can invoke the cast () function.



### API cast

```
out = df.select(
    pl.col("integers").cast(pl.Float32).alias("integers_as_floats"),
    pl.col("floats").cast(pl.Int32).alias("floats_as_integers"),
    pl.col("floats_with_decimal")
    .cast(pl.Int32)
    .alias("floats_with_decimal_as_integers"),
)
print(out)
```

```
shape: (5, 3)

integers as floats | floats as integers | floats with decimal as integers | --- | --- | --- | | f32 | i32 | i32 | i33 | i33 | i34 | i35 |
```

Note that in the case of decimal values these are rounded downwards when casting to an integer.

### Downcast

Reducing the memory footprint is also achievable by modifying the number of bits allocated to an element. As an illustration, the code below demonstrates how casting from Int64 to Int16 and from Float64 to Float32 can be used to lower memory usage.



### API cast

```
out = df.select(
    pl.col("integers").cast(pl.Int16).alias("integers_smallfootprint"),
    pl.col("floats").cast(pl.Float32).alias("floats_smallfootprint"),
)
print(out)
```

# 

### Overflow

When performing downcasting, it is crucial to ensure that the chosen number of bits (such as 64, 32, or 16) is sufficient to accommodate the largest and smallest numbers in the column. For example, using a 32-bit signed integer (Int32) allows handling integers within the range of -2147483648 to +2147483647, while using Int8 covers integers between -128 to 127. Attempting to cast to a DataType that is too small will result in a ComputeError thrown by Polars, as the operation is not supported.



### API cast

```
try:
    out = df.select(pl.col("big_integers").cast(pl.Int8))
    print(out)
except Exception as e:
    print(e)
```

strict conversion from `i64` to `i8` failed for value(s) [10000002, 10000004, 10000005]; if you were trying to cast Utf8 to temporal dtypes, consider using `strptime`

You can set the strict parameter to False, this converts values that are overflowing to null values.



### API cast

```
out = df.select(pl.col("big_integers").cast(pl.Int0, strict=False))
print(out)
```

#### **Strings**

Strings can be casted to numerical data types and vice versa:



### API cast

In case the column contains a non-numerical value, Polars will throw a ComputeError detailing the conversion error. Setting strict=False will convert the non float value to null.



### API cast

```
df = pl.DataFrame({"strings_not_float": ["4.0", "not_a_number", "6.0", "7.0", "8.0"]})
try:
    out = df.select(pl.col("strings_not_float").cast(pl.Float64))
    print(out)
except Exception as e:
    print(e)
```

```
strict conversion from `str` to `f64` failed for value(s) ["not_a_number"]; if you were trying to cast Utf8 to temporal dtypes, consider using `strptime`
```

### **Booleans**

Booleans can be expressed as either 1 (True) or 0 (False). It's possible to perform casting operations between a numerical DataType and a boolean, and vice versa. However, keep in mind that casting from a string (Utf8) to a boolean is not permitted.



### API cast

```
integers | floats | --- | --- | bool | bool | true | false | false | true | tru
```

### **Dates**

Temporal data types such as <code>Date</code> or <code>Datetime</code> are represented as the number of days (<code>Date</code>) and microseconds (<code>Datetime</code>) since epoch. Therefore, casting between the numerical types and the temporal data types is allowed.



### API cast

To perform casting operations between strings and Dates / Datetimes, strftime and strptime are utilized. Polars adopts the chrono format syntax for when formatting. It's worth noting that strptime features additional options that support timezone functionality. Refer to the API documentation for further information.



### API strftime API strptime

### 4.4.5 Strings

The following section discusses operations performed on Utf8 strings, which are a frequently used DataType when working with DataFrames. However, processing strings can often be inefficient due to their unpredictable memory size, causing the CPU to access many random memory locations. To address this issue, Polars utilizes Arrow as its backend, which stores all strings in a contiguous block of memory. As a result, string traversal is cache-optimal and predictable for the CPU.

String processing functions are available in the str namespace.

#### Accessing the string namespace

The str namespace can be accessed through the .str attribute of a column with Utf8 data type. In the following example, we create a column named animal and compute the length of each element in the column in terms of the number of bytes and the number of characters. If you are working with ASCII text, then the results of these two computations will be the same, and using lengths is recommended since it is faster.



### API lengths API n\_chars

```
df = pl.DataFrame({"animal": ["Crab", "cat and dog", "rab$bit", None]})

out = df.select(
    pl.col("animal").str.lengths().alias("byte_count"),
    pl.col("animal").str.n_chars().alias("letter_count"),
)
print(out)
```

```
byte_count | letter_count | --- | --- | u32 | u32 | u32 | u31 | u32 | u32 | u32 | u31 | u32 | u31 | u3
```

# String Parsing

Polars offers multiple methods for checking and parsing elements of a string. Firstly, we can use the contains method to check whether a given pattern exists within a substring. Subsequently, we can extract these patterns and replace them using other methods, which will be demonstrated in upcoming examples.

### Check for existence of a pattern

To check for the presence of a pattern within a string, we can use the contains method. The contains method accepts either a regular substring or a regex pattern, depending on the value of the literal parameter. If the pattern we're searching for is a simple substring located either at the beginning or end of the string, we can alternatively use the starts\_with and ends\_with functions.



### $\label{eq:contains} \textbf{API} \ \ \text{starts\_with} \ \ \textbf{API} \ \ \text{ends\_with}$

```
out = df.select(
    pl.col("animal"),
    pl.col("animal").str.contains("cat|bit").alias("regex"),
    pl.col("animal").str.contains("rab$", literal=True).alias("literal"),
    pl.col("animal").str.starts_with("rab").alias("starts_with"),
    pl.col("animal").str.ends_with("dog").alias("ends_with"),
}
print(out)
```

```
shape: (4, 5)

animal | regex | literal | starts_with | ends_with | --- | --- | --- | --- | --- | str | bool | bool | bool | bool | Crab | false | false | false | false |
```

```
cat and dog | true | false | true | rab$bit | true | true | false | null | null | null | null |
```

### Extract a pattern

The extract method allows us to extract a pattern from a specified string. This method takes a regex pattern containing one or more capture groups, which are defined by parentheses () in the pattern. The group index indicates which capture group to output.



### API extract

```
shape: (3, 1)

a
---
str
messi
null
ronaldo
```

To extract all occurrences of a pattern within a string, we can use the <code>extract\_all</code> method. In the example below, we extract all numbers from a string using the regex pattern (\d+), which matches one or more digits. The resulting output of the <code>extract\_all</code> method is a list containing all instances of the matched pattern within the string.

```
Python
```

# API extract\_all

```
df = pl.DataFrame(("foo": ["123 bla 45 asd", "xyz 678 910t"]))
out = df.select(
    pl.col("foo").str.extract_all(r"(\d+)").alias("extracted_nrs"),
)
print(out)
```

```
shape: (2, 1)

extracted_nrs
---
list[str]

["123", "45"]
["678", "910"]
```

### Replace a pattern

We have discussed two methods for pattern matching and extraction thus far, and now we will explore how to replace a pattern within a string. Similar to extract and extract\_all, Polars provides the replace and replace\_all methods for this purpose. In the example below we replace one match of abc at the end of a word ( \b ) by ABC and we replace all occurrence of a with -.



```
API replace API replace_all
```

```
df = pl.DataFrame({"id": [1, 2], "text": ["123abc", "abc456"]})
out = df.with_columns(
    pl.col("text").str.replace(r"abc\b", "ABC"),
    pl.col("text").str.replace_all("a", "-", literal=True).alias("text_replace_all"),
)
print(out)
```

### **API Documentation**

In addition to the examples covered above, Polars offers various other string manipulation methods for tasks such as formatting, stripping, splitting, and more. To explore these additional methods, you can go to the API documentation of your chosen programming language for Polars.

### 4.4.6 Aggregation

Polars implements a powerful syntax defined not only in its lazy API, but also in its eager API. Let's take a look at what that means.

We can start with the simple US congress dataset.





API DataFrame API Categorical

```
url = "https://theunitedstates.io/congress-legislators/legislators-historical.csv"

dtypes = {
    "first_name": pl.Categorical,
    "gender": pl.Categorical,
    "type": pl.Categorical,
    "state": pl.Categorical,
    "party": pl.Categorical,
}

dataset = pl.read_csv(url, dtypes=dtypes).with_columns(
    pl.col("birthday").str.strptime(pl.Date, strict=False)
}
```

API DataFrame API Categorical Available on feature dtype-categorical

### Basic aggregations

You can easily combine different aggregations by adding multiple expressions in a list. There is no upper bound on the number of aggregations you can do, and you can make any combination you want. In the snippet below we do the following aggregations:

 $Per\ GROUP\ "\texttt{first\_name}"\ we$ 

- count the number of rows in the group:
- short form: pl.count("party")
- full form: pl.col("party").count()
- aggregate the gender values groups:
- full form: pl.col("gender")
- $\bullet$  get the first value of column "last\_name" in the group:
- short form: pl.first("last\_name") (not available in Rust)
- full form: pl.col("last name").first()

Besides the aggregation, we immediately sort the result and limit to the top 5 so that we have a nice summary overview.





### API groupby

# API groupby

```
shape: (5, 4)
  first_name
                       count
                                   gender
                                                                  last_name
                                   list[cat]
                                                                  str
                       u32
 cat
                                   ["M", "M", ... "M"]
["M", "M", ... "M"]
["M", "M", ... "M"]
["M", "M", ... "M"]
["M", "M", ... "M"]
  John
                                                                  Walker
  William
                                                                  Few
                                                                  Armstrong
  Thomas
                       454
439
                                                                  Tucker
  Charles
                                                                Carroll
```

### Conditionals

It's that easy! Let's turn it up a notch. Let's say we want to know how many delegates of a "state" are "Pro" or "Anti" administration. We could directly query that in the aggregation without the need of a lambda or grooming the DataFrame.





### API groupby

### API groupby

Similarly, this could also be done with a nested GROUPBY, but that doesn't help show off some of these nice features.





### API groupby

### API groupby

```
shape: (5, 3)
 state
          party
 cat
          cat
                                u32
 VA
          Anti-Administration
 NJ
          Pro-Administration
 CT
          Pro-Administration
          Pro-Administration
 SC
         Pro-Administration
                                1
```

### Filtering

We can also filter the groups. Let's say we want to compute a mean per group, but we don't want to include all values from that group, and we also don't want to filter the rows from the DataFrame (because we need those rows for another aggregation).

In the example below we show how this can be done.



Note that we can make Python functions for clarity. These functions don't cost us anything. That is because we only create Polars expressions, we don't apply a custom function over a Series during runtime of the query. Of course, you can make functions that return expressions in Rust, too.





### API groupby

### API groupby

```
shape: (5, 5)
 state
          avg M birthday
                           avg F birthday
                                                      # female
 cat
          181.593407
                                             97
                                                      0
 DE
                           null
 ND
 TN
          175.949091
                           109.6
                                             297
                                                      5
          139.0
                           67.666667
                                             52
 NM
 KS
          148.397059
                           85.714286
                                             136
```

### Sorting

It's common to see a DataFrame being sorted for the sole purpose of managing the ordering during a GROUPBY operation. Let's say that we want to get the names of the oldest and youngest politicians per state. We could SORT and GROUPBY.





# API groupby

```
def get_person() -> pl.Expr:
    return pl.col("first_name") + pl.lit(" ") + pl.col("last_name")

q = (
    dataset.lazy()
    .sort("birthday", descending=True)
    .groupby("state")
    .agg(
        get_person().first().alias("youngest"),
        get_person().last().alias("oldest"),
    )
    .limit(5)
)

df = q.collect()
print(df)
```

# API groupby

```
fn get_person() -> Expr {
    col("first_name") + lit(" ") + col("last_name")
}

let df = dataset
    .clone()
    .lazy()
    .sort(
        "birthday",
        SortOptions {
            descending: true,
            nulls_last: true,
            },
        },
        .groupby(["state"])
        .agg([
            get_person().first().alias("youngest"),
            get_person().last().alias("oldest"),
])
        .limit(5)
        .collect()?;
println!("()", df);
```

```
shape: (5, 3)
 state
                                 oldest
          youngest
 cat
          str
                                 str
 KS
          Steven Watkins
                                 James Lane
          Xochitl Torres Small
 NM
                                 José Gallegos
 MO
          Vicky Hartzler
                                 Spencer Pettis
 DE
          John Carney
                                 Samuel White
          Stephen Fincher
 TN
                                 William Cocke
```

However, if we also want to sort the names alphabetically, this breaks. Luckily we can sort in a groupby context separate from the DataFrame .





# API groupby

```
def get_person() -> pl.Expr:
    return pl.col("first_name") + pl.lit(" ") + pl.col("last_name")

q = (
    dataset.lazy()
    .sort("birthday", descending=True)
    .groupby("state")
    .agg(
        get_person().first().alias("youngest"),
        get_person().last().alias("oldest"),
        get_person().sort().first().alias("alphabetical_first"),
    )
    .limit(5)
)

df = q.collect()
print(df)
```

# API groupby

```
let df = dataset
    .clone()
    .lazy()
    .sort(
        "birthday",
        SortOptions {
            descending: true,
            nulls_last: true,
        },
     }
     .groupby(["state"])
     .agg([
            get_person().first().alias("youngest"),
            get_person().last().alias("oldest"),
            get_person().sort(false).first().alias("alphabetical_first"),
     ])
     .limit(5)
     .collect()?;
println!("{}", df);
```

```
shape: (5, 4)
 state
                                   oldest
                                                     alphabetical_first
           youngest
                                                     str
 cat
           str
                                   str
 ND
           Rick Berg
                                   Lyman Casey
                                                     Arthur Link
          Xochitl Torres Small
Steven Watkins
                                   José Gallegos
 NM
                                                     Albert Fall
 KS
                                   James Lane
                                                     Abel Wilder
 TN
          Stephen Fincher
                                   William Cocke
                                                     Aaron Brown
 DE
          John Carney
                                   Samuel White
                                                     Albert Polk
```

We can even sort by another column in the groupby context. If we want to know if the alphabetically sorted name is male or female we could add: pl.col("gender").sort by("first name").first().alias("gender")





### API groupby

```
def get_person() -> pl.Expr:
    return pl.col("first_name") + pl.lit(" ") + pl.col("last_name")

q = (
    dataset.lazy()
    .sort("birthday", descending=True)
    .groupby("state")
    .agg(
        get_person().first().alias("youngest"),
        get_person().last().alias("oldest"),
        get_person().sort().first().alias("alphabetical_first"),
        pl.col("gender").sort_by("first_name").first().alias("gender"),
    )
    .sort("state")
    .limit(5)
)

df = q.collect()
print(df)
```

#### API groupby

```
let df = dataset
    .clone()
     .lazy()
     .sort(
        "birthday",
        SortOptions {
             descending: true,
             nulls_last: true,
     .groupby(["state"])
     .agg([
        get_person().first().alias("youngest"),
        get_person().last().alias("oldest"),
get_person().sort(false).first().alias("alphabetical_first"),
         col("gender")
             .sort_by(["first_name"], [false])
.first()
              .alias("gender"),
    .sort("state", SortOptions::default())
    .limit(5)
     .collect()?;
println!("{}", df);
```

state	youngest	oldest	alphabetical_first	gende
cat	str	str	str	cat
	<del></del>			
PA	Conor Lamb	Thomas Fitzsimons	Aaron Kreider	M
OH	Anthony Gonzalez	John Smith	Aaron Harlan	M
VA	Scott Taylor	William Grayson	A. McEachin	M
MA	Joseph Kennedy	William Widgery	Aaron Hobart	M
FL	Patrick Murphy	Charles Downing	Abijah Gilbert	M
i .	1	1	1	

## DO NOT KILL PARALLELIZATION



The following section is specific to Python , and doesn't apply to Rust . Within Rust , blocks and closures (lambdas) can, and will, be executed concurrently.

We have all heard that Python is slow, and does "not scale." Besides the overhead of running "slow" bytecode, Python has to remain within the constraints of the Global Interpreter Lock (GIL). This means that if you were to use a lambda or a custom Python function to apply during a parallelized phase, Polars speed is capped running Python code preventing any multiple threads from executing the function.

This all feels terribly limiting, especially because we often need those lambda functions in a .groupby() step, for example. This approach is still supported by Polars, but keeping in mind bytecode and the GIL costs have to be paid. It is recommended to try to solve your queries using the expression syntax before moving to lambdas. If you want to learn more about using lambdas, go to the user defined functions section.

### CONCLUSION

In the examples above we've seen that we can do a lot by combining expressions. By doing so we delay the use of custom Python functions that slow down the queries (by the slow nature of Python AND the GIL).

If we are missing a type expression let us know by opening a feature request!

# 4.4.7 Missing data

This page sets out how missing data is represented in Polars and how missing data can be filled.

#### null and NaN values

Each column in a DataFrame (or equivalently a Series) is an Arrow array or a collection of Arrow arrays based on the Apache Arrow format. Missing data is represented in Arrow and Polars with a null value. This null missing value applies for all data types including numerical values.

Polars also allows NotaNumber or NaN values for float columns. These NaN values are considered to be a type of floating point data rather than missing data. We discuss NaN values separately below.

You can manually define a missing value with the python None value:



### API DataFrame



In Pandas the value for missing data depends on the dtype of the column. In Polars missing data is always represented as a null value.

### Missing data metadata

Each Arrow array used by Polars stores two kinds of metadata related to missing data. This metadata allows Polars to quickly show how many missing values there are and which values are missing.

The first piece of metadata is the null count - this is the number of rows with null values in the column:



## API null\_count

```
null_count_df = df.null_count()
print(null_count_df)
```

```
shape: (1, 1)

value
---
u32

1
```

The null\_count method can be called on a DataFrame, a column from a DataFrame or a Series. The null\_count method is a cheap operation as null count is already calculated for the underlying Arrow array.

The second piece of metadata is an array called a *validity bitmap* that indicates whether each data value is valid or missing. The validity bitmap is memory efficient as it is bit encoded - each value is either a 0 or a 1. This bit encoding means the memory overhead per array is only (array length / 8) bytes. The validity bitmap is used by the <code>is\_null</code> method in <code>Polars</code>.

You can return a Series based on the validity bitmap for a column in a DataFrame or a Series with the is\_null method:



### API is\_null

```
is_null_series = df.select(
    pl.col("value").is_null(),
)
print(is_null_series)
```

The is\_null method is a cheap operation that does not require scanning the full column for null values. This is because the validity bitmap already exists and can be returned as a Boolean array.

### Filling missing data

Missing data in a Series can be filled with the fill\_null method. You have to specify how you want the fill\_null method to fill the missing data. The main ways to do this are filling with:

- a literal such as 0 or "0"
- a strategy such as filling forwards
- an expression such as replacing with values from another column
- interpolation

We illustrate each way to fill nulls by defining a simple  $\mathtt{DataFrame}$  with a missing value in  $\mathtt{col2}$ :



### API DataFrame

### FILL WITH SPECIFIED LITERAL VALUE

We can fill the missing data with a specified literal value with pl.lit:

```
Python
```

# API fill\_null

### FILL WITH A STRATEGY

We can fill the missing data with a strategy such as filling forward:



### API fill\_null

```
fill_forward_df = df.with_columns(
    pl.col("col2").fill_null(strategy="forward"),
)
print(fill_forward_df)
```

You can find other fill strategies in the API docs.

# FILL WITH AN EXPRESSION

For more flexibility we can fill the missing data with an expression. For example, to fill nulls with the median value from that column:



# API fill\_null

```
fill_median_df = df.with_columns(
    pl.col("col2").fill_null(pl.median("col2")),
)
print(fill_median_df)
```

In this case the column is cast from integer to float because the median is a float statistic.

# FILL WITH INTERPOLATION

In addition, we can fill nulls with interpolation (without using the fill\_null function):



# API interpolate

```
fill_interpolation_df = df.with_columns(
    pl.col("col2").interpolate(),
)
print(fill_interpolation_df)
```

### NotaNumber or NaN values

Missing data in a Series has a null value. However, you can use NotaNumber or NaN values in columns with float datatypes. These NaN values can be created from Numpy's np.nan or the native python float('nan'):



# API DataFrame



In Pandas by default a Nan value in an integer column causes the column to be cast to float. This does not happen in Polars - instead an exception is raised.

Nan values are considered to be a type of floating point data and are not considered to be missing data in Polars. This means:

- NaN values are not counted with the null\_count method
- Nan values are filled when you use fill nan method but are not filled with the fill null method


Polars has is\_nan and fill\_nan methods which work in a similar way to the is\_null and fill\_null methods. The underlying Arrow arrays do not have a pre-computed validity bitmask for NaN values so this has to be computed for the is\_nan method.

One further difference between null and NaN values is that taking the mean of a column with null values excludes the null values from the calculation but with NaN values taking the mean results in a NaN. This behaviour can be avoided by replacing the NaN values with null values;



# API fill\_nan

```
mean_nan_df = nan_df.with_columns(
    pl.col("value").fill_nan(None).alias("value"),
).mean()
print(mean_nan_df)
```



# 4.4.8 Window functions

Window functions are expressions with superpowers. They allow you to perform aggregations on groups in the select context. Let's get a feel for what that means. First we create a dataset. The dataset loaded in the snippet below contains information about pokemon:





# API read\_csv

```
import polars as pl

# then let's load some csv data with information about pokemon
df = pl.read_csv(
    "https://gist.githubusercontent.com/ritchie46/cac6b337ea52281aa23c049250a4ff03/raw/89a957ff3919d90e6ef2d34235e6bf22304f3366/pokemon.csv"
)
print(df.head())
```

# API CsvReader · Available on feature csv

```
use polars::prelude::*;
use request::blocking::Client;

let data: Vec<u8> = Client::new()
    .get("https://gist.githubusercontent.com/ritchie46/cac6b337ea52281aa23c049250a4ff03/raw/89a957ff3919d90e6ef2d34235e6bf22304f3366/pokemon.csv")
    .send()?
    .text()?
    .bytes()
    .collect();

let df = CsvReader::new(std::io::Cursor::new(data))
    .has_header(true)
    .finish()?;

println!("{}", df);
```

shape:	(5, 13)							
#	Name	Type 1	Type :	2	.   Sp.	Def Spee	d Gener	ation   Legendary
i64	str	str	str		i64	i64	i64	bool
<del></del>				-		+		<del></del>
1	Bulbasaur	Grass	Poison	ı	. 65	45	1	false
2	Ivysaur	Grass	Poison	ı	. 80	60	1	false
3	Venusaur	Grass	Poison	ı	.   100	80	1	false
3	VenusaurMega Venusaur	Grass	Poison	ı	.   120	80	1	false
4	Charmander	Fire	null		.   50	65	1	false
<u>i</u>	1							

### **Groupby Aggregations in selection**

Below we show how to use window functions to group over different columns and perform an aggregation on them. Doing so allows us to use multiple groupby operations in parallel, using a single query. The results of the aggregation are projected back to the original rows. Therefore, a window function will almost always lead to a DataFrame with the same size as the original.

We will discuss later the cases where a window function can change the numbers of rows in a  $\mbox{\tt DataFrame}$  .

Note how we call .over("Type 1") and .over(["Type 1", "Type 2"]) . Using window functions we can aggregate over different groups in a single select call! Note that, in Rust, the type of the argument to over() must be a collection, so even when you're only using one column, you must provided it in an array.

The best part is, this won't cost you anything. The computed groups are cached and shared between different window expressions.





# API over

```
out = df.select(
   "Type 1",
   "Type 2",
   pl.col("Attack").mean().over("Type 1").alias("avg_attack_by_type"),
   pl.col("Defense")
   .mean()
   .over(["Type 1", "Type 2"])
   .alias("avg_defense_by_type_combination"),
   pl.col("Attack").mean().alias("avg_attack"),
)
print(out)
```

### API over

'				
Type 1	Type 2	avg_attack_by_type		
str	str	£64	f64	f64
Grass	Poison	72.923077	67.8	75.349693
Grass	Poison	72.923077	67.8	75.349693
Grass	Poison	72.923077	67.8	75.349693
Grass	Poison	72.923077	67.8	75.349693
Dragon	null	94.0	55.0	75.349693
Dragon	null	94.0	55.0	75.349693
Dragon	Flying	94.0	95.0	75.349693
Psychic	null	53.875	51.428571	75.349693

# Operations per group

Window functions can do more than aggregation. They can also be viewed as an operation within a group. If, for instance, you want to sort the values within a group, you can write col("value").sort().over("group") and voilà! We sorted by group!

Let's filter out some rows to make this more clear.





# API filter

```
filtered = df.filter(pl.col("Type 2") == "Psychic").select(
    "Name",
    "Type 1",
    "Speed",
)
print(filtered)
```

# API filter

```
let filtered = df
   .clone()
   .lazy()
   .filter(col("Type 2").eq(lit("Psychic")))
   .select([col("Name"), col("Type 1"), col("Speed")])
   .collect()?;
println!("{}", filtered);
```

```
shape: (7, 3)
 Name
                        Type 1
                                  Speed
                                  i64
 str
                        str
 Slowpoke
                        Water
 Slowbro
                        Water
 SlowbroMega Slowbro
                        Water
 Exeggcute
                        Grass
                                  40
55
 Exeggutor
 Starmie
                        Water
 Jynx
                        Ice
                                  95
```

Observe that the group Water of column Type 1 is not contiguous. There are two rows of Grass in between. Also note that each pokemon within a group are sorted by Speed in ascending order. Unfortunately, for this example we want them sorted in descending speed order. Luckily with window functions this is easy to accomplish.





# API over

```
out = filtered.with_columns(
    pl.col(["Name", "Speed"]).sort_by("Speed", descending=True).over("Type 1"),
)
print(out)
```

# API over

```
let out = filtered
   .lazy()
   .with_columns([cols(["Name", "Speed"]).sort_by(["Speed"],[true]).over(["Type 1"])])
   .collect()?;
println!("{}", out);
```

```
shape: (7, 3)
Name
                                  Speed
                        Type 1
 str
                         str
                                  i64
                                  115
 Starmie
                        Water
                        Water
                                  30
 SlowbroMega Slowbro
                        Water
                        Grass
 Exeggutor
                        Grass
                                  40
 Slowpoke
                        Water
 Jynx
                         Ice
```

Polars keeps track of each group's location and maps the expressions to the proper row locations. This will also work over different groups in a single select.

The power of window expressions is that you often don't need a groupby -> explode combination, but you can put the logic in a single expression. It also makes the API cleaner. If properly used a:

- groupby -> marks that groups are aggregated and we expect a DataFrame of size n\_groups
- · over -> marks that we want to compute something within a group, and doesn't modify the original size of the DataFrame except in specific cases

### Map the expression result to the DataFrame rows

In cases where the expression results in multiple values per group, the Window function has 3 strategies for linking the values back to the DataFrame rows:

- mapping strategy = 'group to rows' -> each value is assigned back to one row. The number of values returned should match the number of rows.
- mapping\_strategy = 'join' -> the values are imploded in a list, and the list is repeated on all rows. This can be memory intensive.
- mapping strategy = 'explode' -> the values are exploded to new rows. This operation changes the number of rows.

#### Window expression rules

The evaluations of window expressions are as follows (assuming we apply it to a pl.Int32 column):





#### API over

```
# aggregate and broadcast within a group
# output type: -> Int32
pl.sum("foo").over("groups")

# sum within a group and multiply with group elements
# output type: -> Int32
(pl.col("x").sum() * pl.col("y")).over("groups")

# sum within a group and multiply with group elements
# and aggregate the group to a list
# output type: -> List(Int32)
(pl.col("x").sum() * pl.col("y")).over("groups", mapping_strategy="join")

# sum within a group and multiply with group elements
# and aggregate the group to a list
# then explode the list to multiple rows

# This is the fastest method to do things over groups when the groups are sorted
(pl.col("x").sum() * pl.col("y")).over("groups", mapping_strategy="explode")
```

## API over

```
// output type: -> i32
sum("foo").over([col("groups")])
// sum within a group and multiply with group elements // output type: -> i32 \,
(col("x").sum() * col("y"))
    .over([col("groups")])
    .alias("x1")
// sum within a group and multiply with group elements
// and aggregate the group to a list
(col("x").sum() * col("v"))
    .list()
    .over([col("groups")])
    .alias("x2")
// note that it will require an explicit `list()` call
/\!/ sum within a group and multiply with group elements /\!/ and aggregate the group to a list
// the flatten call explodes that list
\ensuremath{//} This is the fastest method to do things over groups when the groups are sorted
(col("x").sum() * col("y"))
    .list()
    .over([col("groups")])
    .flatten()
    .alias("x3");
```

#### More examples

For more exercise, below are some window functions for us to compute:

- sort all pokemon by type
- select the first 3 pokemon per type as "Type 1"
- sort the pokemon within a type by speed in descending order and select the first 3 as "fastest/group"
- sort the pokemon within a type by attack in descending order and select the first 3 as "strongest/group"
- sort the pokemon within a type by name and select the first 3 as "sorted\_by\_alphabet"





### API over API implode

```
out = df.sort("Type 1").select(
    pl.col("Type 1").head(3.over("Type 1", mapping_strategy="explode"),
    pl.col("Name")
    .sort_by(pl.col("Speed"), descending=True)
    .head(3)
    .over("Type 1", mapping_strategy="explode")
    .alias("fastest/group"),
    pl.col("Name")
    .sort_by(pl.col("Attack"), descending=True)
    .head(3)
    .over("Type 1", mapping_strategy="explode")
    .alias("strongest/group"),
    pl.col("Name")
    .sort()
    .head(3)
    .over("Type 1", mapping_strategy="explode")
    .alias("sorted_by_alphabet"),
    )
    print(out)
```

# API over API implode

```
let out = df
     .clone()
     .lazy()
     .select([
         col("Type 1")
             .head(Some(3))
.list()
              .over(["Type 1"])
              .flatten(),
         col("Name")
             .sort_by(["Speed"], [true])
              .head(Some(3))
             .list()
.over(["Type 1"])
              .flatten()
              .alias("fastest/group"),
         col("Name")
             .sort_by(["Attack"], [true])
              .list()
              .over(["Type 1"])
              .flatten()
.alias("strongest/group"),
              .sort(false)
             .head(Some(3))
              .list()
              .over(["Type 1"])
             .flatten()
              .alias("sorted_by_alphabet"),
.collect()?;
println!("{:?}", out);
```

```
shape: (43, 4)
           fastest/group
                                    strongest/group
                                                             sorted_by_alphabet
 Type 1
 str
           str
                                                             str
 Bug
           BeedrillMega Beedrill
                                    PinsirMega Pinsir
                                                             Beedrill
                                    BeedrillMega Beedrill
                                                             BeedrillMega Beedrill
 Bug
           Scyther
           PinsirMega Pinsir
 Bug
                                                             Butterfree
 Dragon
           Dragonite
                                    Dragonite
                                                             Dragonair
           Kabutops
                                    Kabutops
```

	Water	Starmie	GyaradosMega Gyarados	Blastoise
ĺ	Water	Tentacruel	Kingler	BlastoiseMega Blastoise
	Water	Poliwag	Gyarados	Cloyster
			1	

# 4.4.9 Folds

Polars provides expressions/methods for horizontal aggregations like sum, min, mean, etc. by setting the argument axis=1. However, when you need a more complex aggregation the default methods Polars supplies may not be sufficient. That's when folds come in handy.

The fold expression operates on columns for maximum speed. It utilizes the data layout very efficiently and often has vectorized execution.

### MANUAL SUM

Let's start with an example by implementing the sum operation ourselves, with a fold.





### API fold

# API fold\_exprs

```
let df = df!(
    "a" => &[1, 2, 3],
    "b" => &[10, 20, 30],
)?;

let out = df
    .lazy()
    .select([fold_exprs(lit(0), |acc, x| Ok(Some(acc + x)), [col("*")]).alias("sum")])
    .collect()?;
println!("{}", out);
```

The snippet above recursively applies the function  $f(acc, x) \rightarrow acc$  to an accumulator acc and a new column x. The function operates on columns individually and can take advantage of cache efficiency and vectorization.

### CONDITIONAL

In the case where you'd want to apply a condition/predicate on all columns in a DataFrame a fold operation can be a very concise way to express this.





# API fold

# API fold\_exprs

In the snippet we filter all rows where **each** column value is > 1.

# FOLDS AND STRING DATA

Folds could be used to concatenate string data. However, due to the materialization of intermediate columns, this operation will have squared complexity.

Therefore, we recommend using the <code>concat\_str</code> expression for this.





API concat\_str

API concat\_str · Available on feature concat\_str

```
let df = df!(
    "a" => &["a", "b", "c"],
    "b" => &[1, 2, 3],
)?;

let out = df
    .lazy()
    .select([concat_str([col("a"), col("b")], "")])
    .collect()?;
println!("{:?}", out);
```

# 4.4.10 Lists and Arrays

Polars has first-class support for List columns: that is, columns where each row is a list of homogenous elements, of varying lengths. Polars also has an Array datatype, which is analogous to numpy's ndarray objects, where the length is identical across rows.

Note: this is different from Python's list object, where the elements can be of any type. Polars can store these within columns, but as a generic object datatype that doesn't have the special list manipulation features that we're about to discuss.

### Powerful List manipulation

Let's say we had the following data from different weather stations across a state. When the weather station is unable to get a result, an error code is recorded instead of the actual temperature at that time.



#### API DataFrame

### CREATING A LIST COLUMN

For the weather DataFrame created above, it's very likely we need to run some analysis on the temperatures that are captured by each station. To make this happen, we need to first be able to get individual temperature measurements. This is done by:



### API str.split

```
out = weather.with_columns(pl.col("temperatures").str.split(" "))
print(out)
```

One way we could go post this would be to convert each temperature measurement into its own row:



### API DataFrame.explode

```
out = weather.with_columns(pl.col("temperatures").str.split(" ")).explode(
    "temperatures"
)
print(out)
```

```
shape: (49, 2)
station
              temperatures
 str
              str
 Station 1
 Station 1
Station 1
 Station 1
             E1
             16
 Station 5
 Station 5
 Station 5
              24
 Station 5
             E1
```

However, in Polars, we often do not need to do this to operate on the List elements.

#### OPERATING ON LIST COLUMNS

Polars provides several standard operations on List columns. If we want the first three measurements, we can do a head(3). The last three can be obtained via a tail(3), or alternately, via slice (negative indexing is supported). We can also identify the number of observations via lengths. Let's see them in action:



# API Expr.List

```
out = weather.with_columns(pl.col("temperatures").str.split(" ")).with_columns(
    pl.col("temperatures").list.head(3).alias("top3"),
    pl.col("temperatures").list.slice(-3, 3).alias("bottom_3"),
    pl.col("temperatures").list.lengths().alias("obs"),
)
print(out)
```

```
shape: (5, 5)
 station
                   temperatures
                                                                                bottom_3
                                                                                                             obs
                                                   top3
 str
                   list[str]
                                                                                 list[strl
 Station 1
                   ["20", "5", ... "20"]
                                                    ["20", "5", "5"]
                                                                                 ["9", "6", "20"]
                                                                                                              10
                                                                                 ["90", "70", "40"]
["12", "10", "22"]
                   ["18", "8", ... "40"]
["19", "24", ... "22"]
["E2", "E0", ... "6"]
                                                   ["18", "8", "16"]
["19", "24", "E9"
  Station 2
 Station 3
                                                    ["E2", "E0", "15"]
                                                                                 ["17", "13", "6"]
["22", "24", "E1"]
  Station 5
                   ["14", "8", ... "E1"]
                                                   ["14", "8", "E0"]
```



If you find references to the arr API on Stackoverflow or other sources, just replace arr with list, this was the old accessor for the List datatype. arr now refers to the newly introduced Array datatype (see below).

### ELEMENT-WISE COMPUTATION WITHIN LIST S

If we need to identify the stations that are giving the most number of errors from the starting DataFrame, we need to:

- 1. Parse the string input as a List of string values (already done).
- 2. Identify those strings that can be converted to numbers.
- 3. Identify the number of non-numeric values (i.e. null values) in the list, by row.
- 4. Rename this output as errors so that we can easily identify the stations.

The third step requires a casting (or alternately, a regex pattern search) operation to be perform on each element of the list. We can do this using by applying the operation on each element by first referencing them in the pl.element() context, and then calling a suitable Polars expression on them. Let's see how:



### API Expr.List API element

```
out = weather.with_columns(
    pl.col("temperatures")
    .str.split(" ")
    .list.eval(pl.element().cast(pl.Int64, strict=False).is_null())
    .list.sum()
    .alias("errors")
)
print(out)
```

```
station | temperatures | errors | --- | str | str | u32 | |

Station 1 | 20 5 5 E1 7 13 19 9 6 20 | 1 |
Station 2 | 18 8 16 11 23 E2 8 E2 E2 E2 90 7... | 4 |
Station 3 | 19 24 E9 16 6 12 10 22 | 1 |
Station 4 | E2 E0 15 7 8 10 E1 24 17 13 6 | 3 |
Station 5 | 14 8 E0 16 22 24 E1 | 2 |
```

What if we chose the regex route (i.e. recognizing the presence of *any* alphabetical character?)

```
Python
```

### API str.contains

```
out = weather.with_columns(
    pl.col("temperatures")
    .str.split(" ")
    .list.eval(pl.element().str.contains("(?i)[a-z]"))
    .list.sum()
    .alias("errors")
)
print(out)
```

```
station | temperatures | errors | --- | str | str | u32 | station 1 | 20 5 5 E1 7 13 19 9 6 20 | 1 | station 2 | 18 8 16 11 23 E2 8 E2 E2 E2 90 7... | 4 | station 3 | 19 24 E9 16 6 12 10 22 | 1 | station 4 | E2 E0 15 7 8 10 E1 24 17 13 6 | 3 | station 5 | 14 8 E0 16 22 24 E1 | 2
```

If you're unfamiliar with the (?i), it's a good time to look at the documentation for the str.contains function in Polars! The rust regex crate provides a lot of additional regex flags that might come in handy.

### **Row-wise computations**

This context is ideal for computing in row orientation.

We can apply **any** Polars operations on the elements of the list with the <code>list.eval</code> (<code>list().eval</code> in Rust) expression! These expressions run entirely on Polars' query engine and can run in parallel, so will be well optimized. Let's say we have another set of weather data across three days, for different stations:



#### API DataFrame

```
shape: (10, 4)
 station
               day_1
                        day_2
                                 day_3
 str
                        i64
                                 i64
 Station 1
               17
                        15
                                 16
 Station 2
 Station 3
                        10
                                 24
 Station 4
                                 24
                        18
                                 19
 Station 7
                                 23
 Station 9
                                 16
                                 10
 Station 10
```

Let's do something interesting, where we calculate the percentage rank of the temperatures by day, measured across stations. Pandas allows you to compute the percentages of the rank values. Polars doesn't provide a special function to do this directly, but because expressions are so versatile we can create our own percentage rank expression for highest temperature. Let's try that!



# API list.eval

```
rank_pct = (pl.element().rank(descending=True) / pl.col("*").count()).round(2)

out = weather_by_day.with_columns(
    # create the list of homogeneous data
    pl.concat_list(pl.all().exclude("station")).alias("all_temps")
).select(
    # select all columns except the intermediate list
    pl.all().exclude("all_temps"),
    # compute the rank by calling `list.eval`
    pl.col("all_temps").list.eval(rank_pct, parallel=True).alias("temps_rank"),
)

print(out)
```

```
shape: (10, 5)
 station
                day_1
                          day_2
                                   day_3
                                             temps_rank
 str
                i64
                          i64
                                   i64
                                             list[f64]
 Station 1
                                             [0.33, 1.0, 0.67]
 Station 2
                                             [0.83, 0.83, 0.33]
 Station 3
                                             [1.0, 0.67, 0.33]
 Station 4
                22
                                   24
                                             [0.67, 1.0, 0.33]
                                   19
                                             [0.33, 1.0, 0.67]
 Station 8
                          21
                                   23
                                             [1.0, 0.67, 0.33]
[1.0, 0.67, 0.33]
                                   16
 Station 9
                                             [0.33, 0.67, 1.0]
  Station 10
```

### Polars Array S

Array s are a new data type that was recently introduced, and are still pretty nascent in features that it offers. The major difference between a List and an Array is that the latter is limited to having the same number of elements per row, while a List can have a variable number of elements. Both still require that each element's data type is the same.

We can define Array columns in this manner:



# API Array

```
array_df = pl.DataFrame(
    [
        pl.Series("Array_1", [[1, 3], [2, 5]]),
        pl.Series("Array_2", [[1, 7, 3], [8, 1, 0]]),
    ],
    schema={"Array_1": pl.Array(2, pl.Int64), "Array_2": pl.Array(3, pl.Int64)},
)
print(array_df)
```

```
Array_1 | Array_2 | --- | array[i64, 2] | array[i64, 3] | [1, 3] | [1, 7, 3] | [2, 5] | [8, 1, 0]
```

Basic operations are available on it:

```
Python
```

### API arr

```
out = array_df.select(
   pl.col("Array_1").arr.min().suffix("_min"),
   pl.col("Array_2").arr.sum().suffix("_sum"),
)
print(out)
```

 $Polars \ \ {\tt Array} \ s \ are \ still \ being \ actively \ developed, \ so \ this \ section \ will \ likely \ change \ in \ the \ future.$ 

# 4.4.11 User Defined functions

You should be convinced by now that polar expressions are so powerful and flexible that there is much less need for custom python functions than in other libraries.

Still, you need to have the power to be able to pass an expression's state to a third party library or apply your black box function over data in polars.

For this we provide the following expressions:

- map
- apply

# To map or to apply.

These functions have an important distinction in how they operate and consequently what data they will pass to the user.

A map passes the Series backed by the expression as is.

map follows the same rules in both the select and the groupby context, this will mean that the Series represents a column in a DataFrame. Note that in the groupby context, that column is not yet aggregated!

Use cases for map are for instance passing the Series in an expression to a third party library. Below we show how we could use map to pass an expression column to a neural network model.





### API map

```
df.with_columns([
    pl.col("features").map(lambda s: MyNeuralNetwork.forward(s.to_numpy())).alias("activations")
])
```

```
df.with_columns([
    col("features").map(|s| Ok(my_nn.forward(s))).alias("activations")
])
```

Use cases for map in the groupby context are slim. They are only used for performance reasons, but can quite easily lead to incorrect results. Let me explain why.





API map

# API map

In the snippet above we groupby the "keys" column. That means we have the following groups:

```
"a" -> [10, 7]
"b" -> [1]
```

```
"a" -> [null, 10]
"b" -> [null]
```

Now, let's print and see what we've got.

print(out)

```
shape: (2, 3)

keys | shift_map | shift_expression |
--- | --- | --- |
str | list[i64] | list[i64]

a | [null, 10] | [null, 10] |
b | [7] | [null]
```

Ouch.. we clearly get the wrong results here. Group "b" even got a value from group "a" .

This went horribly wrong, because the map applies the function before we aggregate! So that means the whole column [10, 7, 1] got shifted to [null, 10, 7] and was then aggregated.

So my advice is to never use map in the groupby context unless you know you need it and know what you are doing.

### To apply

Luckily we can fix previous example with apply. apply works on the smallest logical elements for that operation.

That is:

- select context -> single elements
- groupby context -> single groups

So with apply we should be able to fix our example:





### API apply

```
out = df.groupby("keys", maintain_order=True).agg(
    pl.col("values").apply(lambda s: s.shift()).alias("shift_map"),
    pl.col("values").shift().alias("shift_expression"),
)
print(out)
```

#### API apply

```
shape: (2, 3)

keys | shift_map | shift_expression |
--- | --- | --- |
str | list[i64] | list[i64]

a | [null, 10] | [null, 10] |
b | [null] | [null]
```

And observe, a valid result!

# ${\tt apply}$ in the ${\tt select}$ context

In the select context, the apply expression passes elements of the column to the python function.

Note that you are now running python, this will be slow.

Let's go through some examples to see what to expect. We will continue with the DataFrame we defined at the start of this section and show an example with the apply function and a counter example where we use the expression API to achieve the same goals.

### ADDING A COUNTER

In this example we create a global counter and then add the integer 1 to the global state at every element processed. Every iteration the result of the increment will be added to the element value.

Note, this example isn't provided in Rust. The reason is that the global counter value would lead to data races when this apply is evaluated in parallel. It would be possible to wrap it in a Mutex to protect the variable, but that would be obscuring the point of the example. This is a case where the Python Global Interpreter Lock's performance tradeoff provides some safety guarantees.





# API apply

```
counter = 0

def add_counter(val: int) -> int:
    global counter
    counter += 1
    return counter + val

out = df.select(
    pl.col("values").apply(add_counter).alias("solution_apply"),
        (pl.col("values") + pl.arange(1, pl.count() + 1)).alias("solution_expr"),
)
print(out)
```

# API apply

```
| solution_apply | solution_expr | --- | i-- | i
```

### COMBINING MULTIPLE COLUMN VALUES

If we want to have access to values of different columns in a single <code>apply</code> function call, we can create <code>struct</code> data type. This data type collects those columns as fields in the <code>struct</code>. So if we'd create a struct from the columns "keys" and "values", we would get the following struct elements:

```
[
    {"keys": "a", "values": 10},
    {"keys": "a", "values": 7},
    {"keys": "b", "values": 1},
]
```

In Python, those would be passed as dict to the calling python function and can thus be indexed by field: str. In rust, you'll get a Series with the Struct type. The fields of the struct can then be indexed and downcast.





API apply API struct

```
out = df.select(
  pl.struct(["keys", "values"])
  .apply(lambda x: len(x["keys"]) + x["values"])
  .alias("solution_apply"),
  (pl.col("keys").str.lengths() + pl.col("values")).alias("solution_expr"),
)
print(out)
```

API apply : API Struct : Available on feature dtype-struct

```
let out = df
      .lazy()
      .select([
            // pack to struct to get access to multiple fields in a custom `apply/map
           as_struct(&[col("keys"), col("values")])

// we will compute the len(a) + b
                  .apply(
                             // downcast to struct
                             let ca = s.struct_()?;
                             // get the fields as Series
                             let s_a = &ca.fields()[0];
let s_b = &ca.fields()[1];
                             // downcast the `Series` to their known type let ca_a = s_a.utf8()?; let ca_b = s_b.i32()?;
                             // iterate both `ChunkedArrays
                                  .into iter()
                                   .zip(ca_b)
                                   .map(|(opt_a, opt_b)| match (opt_a, opt_b) {
    (Some(a), Some(b)) => Some(a.len() as i32 + b),
                                        _ => None,
                                   .collect();
                             Ok(out.into series())
                       GetOutput::from_type(DataType::Int32),
                  .alias("solution_apply"),
            (\texttt{col}\,(\texttt{"keys"})\,.\texttt{str}\,()\,.\texttt{count\_match}\,(\texttt{"."}) \,\,+\,\,\texttt{col}\,(\texttt{"values"})\,)\,.\texttt{alias}\,(\texttt{"solution\_expr"})\,,
.collect()?;
println!("{}", out);
```

Structs are covered in detail in the next section.

### RETURN TYPES?

Custom python functions are black boxes for polars. We really don't know what kind of black arts you are doing, so we have to infer and try our best to understand what you meant.

As a user it helps to understand what we do to better utilize custom functions.

The data type is automatically inferred. We do that by waiting for the first non-null value. That value will then be used to determine the type of the Series.

# The mapping of python types to polars data types is as follows:

- int -> Int64
- float -> Float64
- bool -> Boolean
- str -> Utf8
- list[tp] -> List[tp] (where the inner type is inferred with the same rules)
- dict[str, [tp]] -> struct
- Any -> object (Prevent this at all times)

# Rust types map as follows:

- i32 or i64 -> Int64
- f32 **or** f64 -> Float64
- bool -> Boolean
- String or str -> Utf8
- Vec < tp > List[tp] (where the inner type is inferred with the same rules)

# 4.4.12 The Struct datatype

Polars Struct s are the idiomatic way of working with multiple columns. It is also a free operation i.e. moving columns into Struct s does not copy any data!

For this section, let's start with a DataFrame that captures the average rating of a few movies across some states in the U.S.:

```
Python
```

### API DataFrame

```
shape: (10, 4)
                     Avg_Rating | Count
 Movie
          Theatre
 str
           str
                                    i64
 Cars
          NE
                      4.5
                                    30
          ME
                     4.4
 ET
           IL
                      4.6
                                    26
 Cars
          ND
                      4.3
                                    29
                                    ...
28
                      4.7
 Cars
          NE
 Up
ET
           IL
                                    33
          SD
                      4.6
                                    26
```

# Encountering the Struct type

A common operation that will lead to a Struct column is the ever so popular value\_counts function that is commonly used in exploratory data analysis. Checking the number of times a state appears the data will be done as so:



# API value\_counts

```
out = ratings.select(pl.col("Theatre").value_counts(sort=True))
print(out)
```

```
Theatre
---
struct[2]

["NE",3]
{"NE",3}
{"IL",3}
{"S",2}
{"ME",1}
{"MD",1}
```

Quite unexpected an output, especially if coming from tools that do not have such a data type. We're not in peril though, to get back to a more familiar output, all we need to do is unnest the Struct column into its constituent columns:



### API unnest

```
out = ratings.select(pl.col("Theatre").value_counts(sort=True)).unnest("Theatre")
print(out)
```

# Why value\_counts returns a Struct

Polars expressions always have a Fn(Series) -> Series signature and Struct is thus the data type that allows us to provide multiple columns as input/ouput of an expression. In other words, all expressions have to return a Series object, and Struct allows us to stay consistent with that requirement.

### Structs as dict s

Polars will interpret a dict sent to the Series constructor as a Struct:



# API Series

# Constructing Series objects

Note that Series here was constructed with the name of the series in the begninng, followed by the values. Providing the latter first is considered an anti-pattern in Polars, and must be avoided.

### EXTRACTING INDIVIDUAL VALUES OF A STRUCT

Let's say that we needed to obtain just the movie value in the Series that we created above. We can use the field method to do so:



# API field

```
out = rating_Series.struct.field("Movie")
print(out)
```

```
shape: (2,)
Series: 'Movie' [str]
[
    "Cars"
    "Toy Story"
]
```

# RENAMING INDIVIDUAL KEYS OF A STRUCT

What if we need to rename individual fields of a Struct column? We first convert the rating\_Series object to a DataFrame so that we can view the changes easily, and then use the rename\_fields method:



### API rename\_fields

```
out = (
    rating_Series.to_frame()
    .select(pl.col("ratings").struct.rename_fields(["Film", "State", "Value"]))
    .unnest("ratings")
)
print(out)
```

# Practical use-cases of Struct columns

### IDENTIFYING DUPLICATE ROWS

Let's get back to the ratings data. We want to identify cases where there are duplicates at a Movie and Theatre level. This is where the Struct datatype shines:



# API is\_duplicated API struct

```
out = ratings.filter(pl.struct("Movie", "Theatre").is_duplicated())
print(out)
```

```
shape: (4, 4)
 Movie
          Theatre
                                  Count
                     Avg_Rating
 str
          str
                     f64
                                  i64
 Cars
          NE
                     4.5
                                  30
 ET
          IL
                     4.6
                                  26
 Cars
          NE
 ET
          IL
                     4.9
                                  26
```

We can identify the unique cases at this level also with <code>is\_unique!</code>

### MULTI-COLUMN RANKING

Suppose, given that we know there are duplicates, we want to choose which rank gets a higher priority. We define *Count* of ratings to be more important than the actual <code>Avg\_Rating</code> themselves, and only use it to break a tie. We can then do:



# API is\_duplicated API struct

```
out = ratings.with_columns(
    pl.struct("Count", "Avg_Rating")
    .rank("dense", descending=True)
    .over("Movie", "Theatre")
    .alias("Rank")
) .filter(pl.struct("Movie", "Theatre").is_duplicated())
print(out)
```

	Т			
Movie	Theatre	Avg_Rating	Count	Rank
str	str	f64	i64	u32
+			+	
Cars	NE	4.5	30	1
ET	IL	4.6	26	2
Cars	NE	4.7	28	2
ET	IL	4.9	26	1

That's a pretty complex set of requirements done very elegantly in Polars!

### USING MULTI-COLUMN APPLY

This was discussed in the previous section on *User Defined Functions*.

# 4.4.13 Numpy

Polars expressions support NumPy ufuncs. See here for a list on all supported numpy functions.

This means that if a function is not provided by Polars, we can use NumPy and we still have fast columnar operation through the NumPy API.

#### EXAMPLE



```
API DataFrame API log · Available on feature numpy
```

```
import polars as pl
import numpy as np

df = pl.DataFrame({"a": [1, 2, 3], "b": [4, 5, 6]})

out = df.select(np.log(pl.all()).suffix("_log"))
print(out)
```

```
shape: (3, 2)

a_log | b_log |
--- | --- |
f64 | f64 |

0.0 | 1.386294 |
0.693147 | 1.609438 |
1.098612 | 1.791759
```

# INTEROPERABILITY

Polars Series have support for NumPy universal functions (ufuncs). Element-wise functions such as np.exp(), np.cos(), np.div(), etc. all work with almost zero overhead.

However, as a Polars-specific remark: missing values are a separate bitmask and are not visible by NumPy. This can lead to a window function or a np.convolve() giving flawed or incomplete results.

Convert a Polars Series to a NumPy array with the .to\_numpy() method. Missing values will be replaced by np.nan during the conversion. If the Series does not include missing values, or those values are not desired anymore, the .view() method can be used instead, providing a zero-copy NumPy array of the data.

# 4.5 Transformations

# 4.5.1 Joins

# Join strategies

 ${\tt Polars}\ \ supports\ the\ following\ join\ strategies\ by\ specifying\ the\ \ {\tt strategy}\ \ argument:$ 

- inner
- left
- outer
- cross
- asof
- semi
- anti

### INNER JOIN

An inner join produces a DataFrame that contains only the rows where the join key exists in both DataFrames. Let's take for example the following two DataFrames:



### API DataFrame

# Python

# API DataFrame

```
shape: (3, 3)

order_id | customer_id | amount |
--- | --- | --- |
str | i64 | i64 |
a | 1 | 100 |
b | 2 | 200 |
c | 2 | 300 |
```

To get a DataFrame with the orders and their associated customer we can do an inner join on the customer\_id column:



### API join

```
df_inner_customer_join = df_customers.join(df_orders, on="customer_id", how="inner")
print(df_inner_customer_join)
```

shape: (3, 4)			
l		1	
customer_id	name	order_ic	d amount
i64	str	str	i64
<u> </u>			+
1	Alice	a	100
2	Bob	b	200
2	Bob	c	300
i ı	1		1 1

#### LEFT JOIN

The left join produces a DataFrame that contains all the rows from the left DataFrame and only the rows from the right DataFrame where the join key exists in the left DataFrame. If we now take the example from above and want to have a DataFrame with all the customers and their associated orders (regardless of whether they have placed an order or not) we can do a left join:



### API join

```
df_left_join = df_customers.join(df_orders, on="customer_id", how="left")
print(df_left_join)
```

```
shape: (4, 4)
 customer_id
                 name
                            order id
                                        amount
 i64
                 str
                            str
                                        i64
 1
                 Alice
                                        100
                                        200
                 Bob
                                        300
                 Charlie
                            null
                                        null
```

Notice, that the fields for the customer with the <code>customer\_id</code> of 3 are null, as there are no orders for this customer.

### OUTER JOIN

The outer join produces a DataFrame that contains all the rows from both DataFrames. Columns are null, if the join key does not exist in the source DataFrame. Doing an outer join on the two DataFrames from above produces a similar DataFrame to the left join:



# API join

```
df_outer_join = df_customers.join(df_orders, on="customer_id", how="outer")
print(df_outer_join)
```

```
shape: (4, 4)
 customer id
                name
                           order id
                                       amount
 i64
                str
                                       i64
                Alice
                                       100
                Bob
                Bob
                                        300
                Charlie
                           null
                                       null
```

#### **CROSS JOIN**

A cross join is a cartesian product of the two DataFrames. This means that every row in the left DataFrame is joined with every row in the right DataFrame. The cross join is useful for creating a DataFrame with all possible combinations of the columns in two DataFrames. Let's take for example the following two DataFrames.



### API DataFrame

```
color color red green green
```

## Python

### API DataFrame

We can now create a DataFrame containing all possible combinations of the colors and sizes with a cross join:



## API join

```
df_cross_join = df_colors.join(df_sizes, how="cross")
print(df_cross_join)
```

```
shape: (9, 2)
 color
          size
 str
          str
         M
L
 red
 red
 blue
         M
 blue
 blue
 green
         M
L
 green
 green
```

The inner, left, outer and cross join strategies are standard amongst dataframe libraries. We provide more details on the less familiar semi, anti and asof join strategies below.

#### SEMI JOIN

Consider the following scenario: a car rental company has a DataFrame showing the cars that it owns with each car having a unique id.



### API DataFrame

```
shape: (3, 2)

id make |
--- | --- |
str | str |
a | ford |
b | toyota |
c | bmw |
```

The company has another DataFrame showing each repair job carried out on a vehicle.

```
Python
```

#### API DataFrame

You want to answer this question: which of the cars have had repairs carried out?

An inner join does not answer this question directly as it produces a DataFrame with multiple rows for each car that has had multiple repair jobs:



## API join

```
df_inner_join = df_cars.join(df_repairs, on="id", how="inner")
print(df_inner_join)
```

```
shape: (2, 3)

id | make | cost |
--- | --- | --- |
str | str | i64 |
c | bmw | 100 |
```

However, a semi join produces a single row for each car that has had a repair job carried out.

```
Python
```

#### API join

```
df_semi_join = df_cars.join(df_repairs, on="id", how="semi")
print(df_semi_join)
```

#### ANTI JOIN

Continuing this example, an alternative question might be: which of the cars have **not** had a repair job carried out? An anti join produces a DataFrame showing all the cars from  $df_{cars}$  where the id is not present in the  $df_{repairs}$  DataFrame.



#### API join

```
df_anti_join = df_cars.join(df_repairs, on="id", how="anti")
print(df_anti_join)
```

#### ASOF JOIN

An asof join is like a left join except that we match on nearest key rather than equal keys. In Polars we can do an asof join with the join method and specifying strategy="asof". However, for more flexibility we can use the join\_asof method.

Consider the following scenario: a stock market broker has a DataFrame called df trades showing transactions it has made for different stocks.



#### API DataFrame

The broker has another DataFrame called df\_quotes showing prices it has quoted for these stocks.

```
Python
```

### API DataFrame

You want to produce a DataFrame showing for each trade the most recent quote provided before the trade. You do this with <code>join\_asof</code> (using the default strategy = "backward"). To avoid joining between trades on one stock with a quote on another you must specify an exact preliminary join on the stock column with <code>by="stock"</code>.



#### API join\_asof

```
df_asof_join = df_trades.join_asof(df_quotes, on="time", by="stock")
print(df_asof_join)
```

If you want to make sure that only quotes within a certain time range are joined to the trades you can specify the tolerance argument. In this case we want to make sure that the last preceding quote is within 1 minute of the trade so we set tolerance = "1m".

```
Python
```

```
df_asof_tolerance_join = df_trades.join_asof(
    df_quotes, on="time", by="stock", tolerance="lm"
)
print(df_asof_tolerance_join)
```

```
shape: (4, 4)

time | stock | trade | quote |
```

datetime[µs]	str	i64	164
2020-01-01 09:01:00	Δ.	101	100
2020-01-01 09:01:00			
2020-01-01 09:03:00			
2020-01-01 09:06:00	C	500	null
L		1	

## 4.5.2 Concatenation

There are a number of ways to concatenate data from separate DataFrames:

- two dataframes with the same columns can be vertically concatenated to make a longer dataframe
- two dataframes with the same number of rows and non-overlapping columns can be horizontally concatenated to make a wider dataframe
- two dataframes with **different numbers of rows and columns** can be **diagonally** concatenated to make a dataframe which might be longer and/or wider. Where column names overlap values will be vertically concatenated. Where column names do not overlap new rows and columns will be added. Missing values will be set as null

### Vertical concatenation - getting longer

In a vertical concatenation you combine all of the rows from a list of  $\mathtt{DataFrames}$  into a single longer  $\mathtt{DataFrames}$ .



### API concat

Vertical concatenation fails when the dataframes do not have the same column names.

### Horizontal concatenation - getting wider

In a horizontal concatenation you combine all of the columns from a list of DataFrames into a single wider DataFrame.



## API concat

Horizontal concatenation fails when dataframes have overlapping columns or a different number of rows.

## Diagonal concatenation - getting longer, wider and null ier

In a diagonal concatenation you combine all of the row and columns from a list of DataFrames into a single longer and/or wider DataFrame.



#### API concat



Diagonal concatenation generates nulls when the column names do not overlap.

When the dataframe shapes do not match and we have an overlapping semantic key then we can join the dataframes instead of concatenating them.

### Rechunking

Before a concatenation we have two dataframes dfl and df2. Each column in dfl and df2 is in one or more chunks in memory. By default, during concatenation the chunks in each column are copied to a single new chunk - this is known as **rechunking**. Rechunking is an expensive operation, but is often worth it because future operations will be faster. If you do not want Polars to rechunk the concatenated <code>DataFrame</code> you specify <code>rechunk = False</code> when doing the concatenation.

## 4.5.3 Pivots

Pivot a column in a DataFrame and perform one of the following aggregations:

- first
- sum
- min
- max
- mean
- median

The pivot operation consists of a group by one, or multiple columns (these will be the new y-axis), the column that will be pivoted (this will be the new x-axis) and an aggregation.

### Dataset



## API DataFrame

## Eager



## API pivot

```
out = df.pivot(index="foo", columns="bar", values="N", aggregate_function="first")
print(out)
```

```
shape: (3, 6)
                               i64
        i64
                i64
                        i64
 str
A
                2
                        null
                                       null
                               null
        null
                null
                                       null
        null
                        null
                               null
                                       2
```

## Lazy

A polars LazyFrame always need to know the schema of a computation statically (before collecting the query). As a pivot's output schema depends on the data, and it is therefore impossible to determine the schema without running the query.

Polars could have abstracted this fact for you just like Spark does, but we don't want you to shoot yourself in the foot with a shotgun. The cost should be clear upfront.



### API pivot

```
q = (
    df.lazy()
    .collect()
    .pivot(index="foo", columns="bar", values="N", aggregate_function="first")
    .lazy()
)
out = q.collect()
print(out)
```

## 4.5.4 Melts

Melt operations unpivot a DataFrame from wide format to long format

## **Dataset**



## API DataFrame

```
shape: (3, 4)

A | B | C | D |

--- | --- | --- |

str | i64 | i64 | i64 |

a | 1 | 10 | 2 |

b | 3 | 11 | 4 |

a | 5 | 12 | 6 |
```

## Eager + Lazy

Eager and lazy have the same API.



### API melt

```
out = df.melt(id_vars=["A", "B"], value_vars=["C", "D"])
print(out)
```

### 4.5.5 Time Series

#### **Parsing**

Polars has native support for parsing time series data and doing more sophisticated operations such as temporal grouping and resampling.

#### DATATYPES

Polars has the following datetime datatypes:

- Date: Date representation e.g. 2014-07-08. It is internally represented as days since UNIX epoch encoded by a 32-bit signed integer.
- Datetime: Datetime representation e.g. 2014-07-08 07:00:00. It is internally represented as a 64 bit integer since the Unix epoch and can have different units such as ns. us. ms.
- Duration : A time delta type that is created when subtracting Date/Datetime . Similar to timedelta in python.
- Time: Time representation, internally represented as nanoseconds since midnight.

#### PARSING DATES FROM A FILE

When loading from a CSV file Polars attempts to parse dates and times if the try\_parse\_dates flag is set to True:



### API read\_csv

```
df = pl.read_csv("docs/src/data/appleStock.csv", try_parse_dates=True)
print(df)
```

On the other hand binary formats such as parquet have a schema that is respected by Polars.

## CASTING STRINGS TO DATES

You can also cast a column of datetimes encoded as strings to a datetime type. You do this by calling the string str.strptime method and passing the format of the date string:



#### API read\_csv API strptime

```
df = pl.read_csv("docs/src/data/appleStock.csv", try_parse_dates=False)

df = df.with_columns(pl.col("Date").str.strptime(pl.Date, format="%Y-%m-%d"))
print(df)
```

```
| 2012-12-04 | 575.85 |
| 2013-07-05 | 417.42 |
| 2013-11-07 | 512.49 |
| 2014-02-25 | 522.06 |
```

The strptime date formats can be found here..

#### EXTRACTING DATE FEATURES FROM A DATE COLUMN

You can extract data features such as the year or day from a date column using the .dt namespace on a date column:



#### API year

```
df_with_year = df.with_columns(pl.col("Date").dt.year().alias("year"))
print(df_with_year)
```

```
shape: (100, 3)
 Date
             Close
                    year
 date
             f64
                      i32
 1981-02-23
             24.62
                     1981
 1981-05-06 27.38
                     1981
1981-05-18
           28.0
                     1981
 1981-09-25
           14.25
 2012-12-04
 2013-07-05
           417.42
                    2013
 2013-11-07
           512.49 | 2013
             522.06 | 2014
 2014-02-25
```

#### MIXED OFFSETS

If you have mixed offsets (say, due to crossing daylight saving time), then you can use utceTrue and then convert to your time zone:



```
API strptime API convert_time_zone · Available on feature timezone
```

```
data = [
    "2021-03-27T00:00:00+0100",
    "2021-03-28T00:00:00+0100",
    "2021-03-29T00:00:00+0200",
    "2021-03-30T00:00:00+0200",
]
mixed_parsed = (
    pl.Series(data)
    .str.strptime(pl.Datetime, format="%Y-%m-%dT%H:%M:%S%z", utc=True)
    .dt.convert_time_zone("Europe/Brussels")
)
print(mixed_parsed)
```

#### **Filtering**

Filtering date columns works in the same way as with other types of columns using the .filter method.

 $Polars\ uses\ Python's\ native\ \texttt{datetime}\ ,\ \texttt{date}\ and\ \texttt{timedelta}\ for\ equality\ comparisons\ between\ the\ \texttt{datatypes}\ \texttt{pl.Datetime}\ ,\ \texttt{pl.Date}\ and\ \texttt{pl.Duration}\ .$ 

In the following example we use a time series of Apple stock prices.

```
Python
```

#### API read\_csv

```
import polars as pl
from datetime import datetime

df = pl.read_csv("docs/src/data/appleStock.csv", try_parse_dates=True)
print(df)
```

```
shape: (100, 2)
 Date
              Close
 date
               f64
 1981-02-23
              24.62
 1981-05-06
 1981-05-18
              28.0
 1981-09-25
              14.25
 2012-12-04
              575.85
             417.42
 2014-02-25
              522.06
```

#### FILTERING BY SINGLE DATES

We can filter by a single date by casting the desired date string to a Date object in a filter expression:



## API filter

```
filtered_df = df.filter(
    pl.col("Date") == datetime(1995, 10, 16),
)
print(filtered_df)
```

```
Date | Close | --- | --- | date | f64 | 1995-10-16 | 36.13 |
```

Note we are using the lowercase datetime method rather than the uppercase Datetime data type.

## FILTERING BY A DATE RANGE

We can filter by a range of dates using the is\_between method in a filter expression with the start and end dates:

```
Python
```

```
API filter API is_between
```

```
filtered_range_df = df.filter(
    pl.col("Date").is_between(datetime(1995, 7, 1), datetime(1995, 11, 1)),
)
print(filtered_range_df)
```

```
Date | Close | --- | date | f64 | 1995-10-16 | 36.13 |
```

## FILTERING WITH NEGATIVE DATES

Say you are working with an archeologist and are dealing in negative dates. Polars can parse and store them just fine, but the Python datetime library does not. So for filtering, you should use attributes in the .dt namespace:



## API strptime

```
ts = pl.Series(["-1300-05-23", "-1400-03-02"]).str.strptime(pl.Date)
negative_dates_df = pl.DataFrame(("ts": ts, "values": [3, 4]))
negative_dates_filtered_df = negative_dates_df.filter(pl.col("ts").dt.year() < -1300)
print(negative_dates_filtered_df)</pre>
```

```
shape: (1, 2)

ts | values | --- | --- | date | i64 | -1400-03-02 | 4
```

### Grouping

### **GROUPING BY FIXED WINDOWS**

We can calculate temporal statistics using groupby\_dynamic to group rows into days/months/years etc.

Annual average example

In following simple example we calculate the annual average closing price of Apple stock prices. We first load the data from CSV:



#### API upsample

```
df = pl.read_csv("docs/src/data/appleStock.csv", try_parse_dates=True)
df = df.sort("Date")
print(df)
```

```
shape: (100, 2)
 Date
               Close
               f64
 1981-02-23
               24.62
  1981-05-06
 1981-05-18
               28.0
 1981-09-25
               14.25
 2012-12-04
               575.85
 2013-07-05
               417.42
 2014-02-25
               522.06
```



 $The \ dates \ are \ sorted \ in \ ascending \ order-if \ they \ are \ not \ sorted \ in \ this \ way \ the \ \ \verb|groupby_dynamic| output \ will \ not \ be \ correct!$ 

To get the annual average closing price we tell <code>groupby\_dynamic</code> that we want to:

- $\bullet$  group by the  ${\tt Date}$  column on an annual (  ${\tt ly}$  ) basis
- $\bullet$  take the mean values of the  ${\tt Close}$  column for each year:



#### API groupby\_dynamic

```
annual_average_df = df.groupby_dynamic("Date", every="1y").agg(pl.col("Close").mean())

df_with_year = annual_average_df.with_columns(pl.col("Date").dt.year().alias("year"))
print(df_with_year)
```

The annual average closing price is then:

```
shape: (34, 3)
 Date
               Close
                            year
               f64
                            i32
 date
 1981-01-01
                            1981
 1982-01-01
               11.0
                            1982
 1983-01-01
               30.543333
 1984-01-01
               27.583333
                            1984
 2011-01-01
                            2011
 2012-01-01
               560.965
 2013-01-01
               464.955
                            2013
 2014-01-01
```

#### Parameters for groupby dynamic

A dynamic window is defined by a:

- every: indicates the interval of the window
- period: indicates the duration of the window
- offset: can be used to offset the start of the windows

The value for every sets how often the groups start. The time period values are flexible - for example we could take:

- the average over 2 year intervals by replacing 1y with 2y
- the average over 18 month periods by replacing 1y with 1y6mo

We can also use the period parameter to set how long the time period for each group is. For example, if we set the every parameter to be 1y and the period parameter to be 2y then we would get groups at one year intervals where each groups spanned two years.

If the period parameter is not specified then it is set equal to the every parameter so that if the every parameter is set to be 1y then each group spans 1y as well.

Because every does not have to be equal to period, we can create many groups in a very flexible way. They may overlap or leave boundaries between them.

Let's see how the windows for some parameter combinations would look. Let's start out boring. 6

- every: 1 day -> "1d"
   period: 1 day -> "1d"
- this creates adjacent windows of the same size

--
--
- every: 1 day -> "1d"
- period: 2 days -> "2d"

```
these windows have an overlap of 1 day
|----|
|----|
|----|
```

- every: 2 days -> "2d"
- period: 1 day -> "1d"

```
this would leave gaps between the windows data points that in these gaps will not be a member of any group |--| |--|
```

#### truncate

The truncate parameter is a Boolean variable that determines what datetime value is associated with each group in the output. In the example above the first data point is on 23rd February 1981. If truncate = True (the default) then the date for the first year in the annual average is 1st January 1981. However, if truncate = False then the date for the first year in the annual average is the date of the first data point on 23rd February 1981. Note that truncate only affects what's shown in the Date column and does not affect the window boundaries.

Using expressions in groupby\_dynamic

We aren't restricted to using simple aggregations like mean in a groupby operation - we can use the full range of expressions available in Polars.

In the snippet below we create a date range with every  $\mathbf{day}$  ( "1d" ) in 2021 and turn this into a DataFrame .

Then in the <code>groupby\_dynamic</code> we create dynamic windows that start every **month** ( "lmo" ) and have a window length of 1 month. The values that match these dynamic windows are then assigned to that group and can be aggregated with the powerful expression API.

Below we show an example where we use groupby\_dynamic to compute:

- the number of days until the end of the month
- the number of days in a month



API groupby\_dynamic .API explode .API date\_range

```
shape: (36, 3)
time
                         day/eom
                                   days_in_month
 \texttt{datetime}\,[\,\mu\,s\,]
 2021-01-01 00:00:00
                         30
                                    31
 2021-01-01 00:00:00
 2021-01-01 00:00:00
                        28
 2021-02-01 00:00:00
                                   28
 2021-11-01 00:00:00
 2021-12-01 00:00:00
                       30
                                   31
 2021-12-01 00:00:00
 2021-12-01 00:00:00
```

### **GROUPING BY ROLLING WINDOWS**

The rolling groupby, <code>groupby\_rolling</code>, is another entrance to the <code>groupby</code> context. But different from the <code>groupby\_dynamic</code> the windows are not fixed by a parameter <code>every</code> and <code>period</code>. In a rolling groupby the windows are not fixed at all! They are determined by the values in the <code>index\_column</code>.

So imagine having a time column with the values {2021-01-06, 2021-01-10} and a period="5d" this would create the following windows:

Because the windows of a rolling groupby are always determined by the values in the DataFrame column, the number of groups is always equal to the original DataFrame.

#### COMBINING GROUPBY'S

Rolling and dynamic groupby's can be combined with normal groupby operations.

Below is an example with a dynamic groupby.



#### API DataFrame

## Python

### API groupby\_dynamic

```
out = df.groupby_dynamic(
    "time",
    every="lh",
    closed="both",
    by="groups",
    include_boundaries=True,
).agg(
    [
        pl.count(),
    ]
)
```

```
shape: (7, 5)
             _lower_boundary
                                      _upper_boundary
                                                               time
                                                                                         count
groups
 str
            \texttt{datetime}\,[\,\mu s\,]
                                      \texttt{datetime}\,[\,\mu s\,]
                                                               \texttt{datetime}[\mu \texttt{s}]
                                                                                         u32
            2021-12-15 23:00:00
                                      2021-12-16 00:00:00
                                                               2021-12-15 23:00:00
 а
 a
a
a
a
b
            2021-12-16 00:00:00
                                      2021-12-16 01:00:00
                                                               2021-12-16 00:00:00
            2021-12-16 01:00:00
                                      2021-12-16 02:00:00
                                                               2021-12-16 01:00:00
            2021-12-16 02:00:00
                                      2021-12-16 03:00:00
                                                               2021-12-16 02:00:00
                                     2021-12-16 04:00:00
2021-12-16 02:00:00
            2021-12-16 03:00:00
                                                               2021-12-16 03:00:00
            2021-12-16 01:00:00
                                                               2021-12-16 01:00:00
            2021-12-16 02:00:00
                                     2021-12-16 03:00:00
                                                             2021-12-16 02:00:00
                                                                                       1
 b
```

#### Resampling

We can resample by either:

- upsampling (moving data to a higher frequency)
- downsampling (moving data to a lower frequency)
- combinations of these e.g. first upsample and then downsample

### DOWNSAMPLING TO A LOWER FREQUENCY

Polars views downsampling as a special case of the **groupby** operation and you can do this with <code>groupby\_dynamic</code> and <code>groupby\_rolling</code> - see the temporal groupby page for examples.

#### **UPSAMPLING TO A HIGHER FREQUENCY**

Let's go through an example where we generate data at 30 minute intervals:



```
API DataFrame 'API date_range
```

```
shape: (7, 3)
time
                                 values
 datetime[µs]
                                  f64
2021-12-16 00:00:00
                                 1.0
 2021-12-16 00:30:00
 2021-12-16 01:00:00
                                 3.0
 2021-12-16 01:30:00
                                 4.0
 2021-12-16 02:00:00
 2021-12-16 02:30:00
                                 6.0
 2021-12-16 03:00:00
                                 7.0
```

Upsampling can be done by defining the new sampling interval. By upsampling we are adding in extra rows where we do not have data. As such upsampling by itself gives a DataFrame with nulls. These nulls can then be filled with a fill strategy or interpolation.

## Upsampling strategies

In this example we upsample from the original 30 minutes to 15 minutes and then use a forward strategy to replace the nulls with the previous non-null value:



#### API upsample

```
out1 = df.upsample(time_column="time", every="15m").fill_null(strategy="forward")
print(out1)
```

In this example we instead fill the nulls by linear interpolation:



```
API upsample API interpolate API fill_null
```

```
out2 = (
    df.upsample(time_column="time", every="15m")
    .interpolate()
    .fill_null(strategy="forward")
)
print(out2)
```

time	groups	values
datetime[µs]	str	f64
	<del> </del>	
2021-12-16 00:00:00	a	1.0
2021-12-16 00:15:00	a	1.5
2021-12-16 00:30:00	a	2.0
2021-12-16 00:45:00	a	2.5
2021-12-16 02:15:00	b	5.5
2021-12-16 02:30:00	a	6.0
2021-12-16 02:45:00	a	6.5
2021-12-16 03:00:00	a	7.0

#### Time zones

## Tom Scott

You really should never, ever deal with time zones if you can help it

The Datetime datatype can have a time zone associated with it. Examples of valid time zones are:

- · None: no time zone, also known as "time zone naive";
- UTC: Coordinated Universal Time;
- · Asia/Kathmandu: time zone in "area/location" format. See the list of tz database time zones to see what's available;
- +01:00 : fixed offsets. May be useful when parsing, but you almost certainly want the "Area/Location" format above instead as it will deal with irregularities such as DST (Daylight Saving Time) for you.

Note that, because a Datetime can only have a single time zone, it is impossible to have a column with multiple time zones. If you are parsing data with multiple offsets, you may want to pass utc=True to convert them all to a common time zone (UTC), see parsing dates and times.

The main methods for setting and converting between time zones are:

- dt.convert\_time\_zone : convert from one time zone to another;
- dt.replace time zone: set/unset/change time zone;

Let's look at some examples of common operations:



API strptime API replace\_time\_zone . Available on feature timezone

```
ts = ["2021-03-27 03:00", "2021-03-28 03:00"]
tz_naive = pl.Series("tz_naive", ts).str.strptime(pl.Datetime)
tz_aware = tz_naive.dt.replace_time_zone("UTC").rename("tz_aware")
time_zones_df = pl.DataFrame([tz_naive, tz_aware])
print(time_zones_df)
```



API convert\_time\_zone API replace\_time\_zone Available on feature timezone

```
shape: (2, 3)

replace time zone | convert time zone | unset time zone | --- | --- | datetime[µs, Europe/Brussels] | datetime[µs, Asia/Kathmandu] | datetime[µs]
```

2021-03-27 03:00:00 CET		2021-03-27 08:45:0	0	+0545		2021-03-27 03:00:00	)
2021-03-28 03:00:00 CEST	1	2021-03-28 08:45:0	0	+0545		2021-03-28 03:00:00	)
	1				1		1

## 4.6 Lazy API

### 4.6.1 Usage

With the lazy API, Polars doesn't run each query line-by-line but instead processes the full query end-to-end. To get the most out of Polars it is important that you use the lazy API because:

- the lazy API allows Polars to apply automatic query optimization with the query optimizer
- the lazy API allows you to work with larger than memory datasets using streaming
- the lazy API can catch schema errors before processing the data

Here we see how to use the lazy API starting from either a file or an existing DataFrame.

### Using the lazy API from a file

In the ideal case we would use the lazy API right from a file as the query optimizer may help us to reduce the amount of data we read from the file.

We create a lazy query from the Reddit CSV data and apply some transformations.

By starting the query with pl.scan csv we are using the lazy API.

```
Python
```

```
API scan_csv API with_columns API filter API col
```

```
q1 = (
    pl.scan_csv(f"docs/src/data/reddit.csv")
    .with_columns(pl.col("name").str.to_uppercase())
    .filter(pl.col("comment_karma") > 0)
)
```

A pl.scan function is available for a number of file types including CSV, IPC, Parquet and JSON.

In this query we tell Polars that we want to:

- load data from the Reddit CSV file
- convert the name column to uppercase
- apply a filter to the comment karma column

The lazy query will not be executed at this point. See this page on executing lazy queries for more on running lazy queries.

## Using the lazy API from a DataFrame

An alternative way to access the lazy API is to call .lazy on a DataFrame that has already been created in memory.



```
API lazy
```

```
q3 = pl.DataFrame({"foo": ["a", "b", "c"], "bar": [0, 1, 2]}).lazy()
```

By calling .lazy we convert the  ${\tt DataFrame}$  to a  ${\tt LazyFrame}$  .

## 4.6.2 Optimizations

If you use Polars 'lazy API, Polars will run several optimizations on your query. Some of them are executed up front, others are determined just in time as the materialized data comes in.

Here is a non-complete overview of optimizations done by polars, what they do and how often they run.

Optimization	Explanation	runs
Predicate pushdown	Applies filters as early as possible/ at scan level.	1 time
Projection pushdown	Select only the columns that are needed at the scan level.	1 time
Slice pushdown	Only load the required slice from the scan level. Don't materialize sliced outputs (e.g. join.head(10)).	1 time
Common subplan elimination	Cache subtrees/file scans that are used by multiple subtrees in the query plan.	1 time
Simplify expressions	Various optimizations, such as constant folding and replacing expensive operations with faster alternatives.	until fixed point
Join ordering	Estimates the branches of joins that should be executed first in order to reduce memory pressure.	1 time
Type coercion	Coerce types such that operations succeed and run on minimal required memory.	until fixed point
Cardinality estimation	Estimates cardinality in order to determine optimal groupby strategy.	0/n times; dependent on query

### 4.6.3 Schema

The schema of a Polars DataFrame or LazyFrame sets out the names of the columns and their datatypes. You can see the schema with the .schema method on a DataFrame or LazyFrame



```
API DataFrame API lazy
```

```
q3 = pl.DataFrame({"foo": ["a", "b", "c"], "bar": [0, 1, 2]}).lazy()
print(q3.schema)
```

```
{'foo': Utf8, 'bar': Int64}
```

The schema plays an important role in the lazy API.

## Type checking in the lazy API

One advantage of the lazy API is that Polars will check the schema before any data is processed. This check happens when you execute your lazy query.

We see how this works in the following simple example where we call the .round expression on the integer bar column.



```
API lazy API with_columns
```

```
pl.DataFrame(("foo": ["a", "b", "c"], "bar": [0, 1, 2])).lazy().with_columns(
    pl.col("bar").round(0)
)
```

The .round expression is only valid for columns with a floating point dtype. Calling .round on an integer column means the operation will raise a SchemaError.

If we executed this query in eager mode the error would only be found once the data had been processed in all earlier steps.

When we execute a lazy query Polars checks for any potential SchemaError before the time-consuming step of actually processing the data in the pipeline.

## The lazy API must know the schema

In the lazy API the Polars query optimizer must be able to infer the schema at every step of a query plan. This means that operations where the schema is not knowable in advance cannot be used with the lazy API.

The classic example of an operation where the schema is not knowable in advance is a .pivot operation. In a .pivot the new column names come from data in one of the columns. As these column names cannot be known in advance a .pivot is not available in the lazy API.

### Dealing with operations not available in the lazy API

If your pipeline includes an operation that is not available in the lazy API it is normally best to:

- run the pipeline in lazy mode up until that point
- execute the pipeline with .collect to materialize a DataFrame
- do the non-lazy operation on the DataFrame
- convert the output back to a LazyFrame with .lazy and continue in lazy mode

We show how to deal with a non-lazy operation in this example where we:

- create a simple DataFrame
- convert it to a LazyFrame with .lazy
- $\bullet$  do a transformation using <code>.with\_columns</code>
- execute the query before the pivot with .collect to get a DataFrame
- do the .pivot on the DataFrame
- · convert back in lazy mode
- do a .filter
- $\bullet$  finish by executing the query with <code>.collect</code> to get a <code>DataFrame</code>



API collect API pivot API filter

```
shape: (2, 4)

| id | jan | feb | mar |
|---| --- | --- |
| str | i64 | i64 | i64 |
| a | 0 | null | null |
| b | null | 2 | null |
```

## 4.6.4 Query Plan

For any lazy query Polars has both:

- a non-optimized plan with the set of steps code as we provided it and
- an optimized plan with changes made by the query optimizer

We can understand both the non-optimized and optimized query plans with visualization and by printing them as text.

Below we consider the following query:

```
Python
```

```
q1 = (
    pl.scan_csv(f"docs/src/data/reddit.csv")
    .with_columns(pl.col("name").str.to_uppercase())
    .filter(pl.col("comment_karma") > 0)
)
```

### Non-optimized query plan

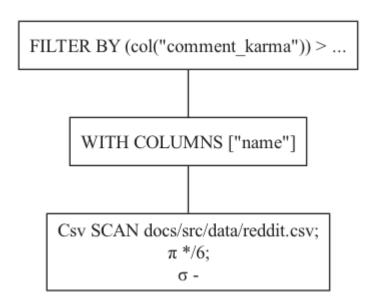
**GRAPHVIZ VISUALIZATION** 

First we visualise the non-optimized plan by setting  ${\tt optimized=False}$ .



API show\_graph

q1.show\_graph(optimized=False)



The query plan visualization should be read from bottom to top. In the visualization:

- each box corresponds to a stage in the query plan
- the sigma stands for SELECTION and indicates any filter conditions
- $\bullet$  the  $\,\mathtt{pi}\,$  stands for  $\,\mathtt{PROJECTION}\,$  and indicates choosing a subset of columns

#### PRINTED QUERY PLAN

We can also print the non-optimized plan with explain(optimized=False)



API explain

q1.explain(optimized=False)

```
FILTER [(col("comment_karma")) > (0)] FROM WITH_COLUMNS:
[col("name").str.uppercase()]

CSV SCAN data/reddit.csv
PROJECT */6 COLUMNS
```

The printed plan should also be read from bottom to top. This non-optimized plan is roughly equal to:

- read from the data/reddit.csv file
- read all 6 columns (where the \* wildcard in PROJECT \*/6 COLUMNS means take all columns)
- transform the name column to uppercase
- apply a filter on the comment\_karma column


#### Optimized query plan

Now we visualize the optimized plan with <code>show\_graph</code>.



API show\_graph

q1.show\_graph()



We can also print the optimized plan with explain



API explain

q1.explain()

```
WITH_COLUMNS:
[col("name").str.uppercase()]

CSV SCAN data/reddit.csv
PROJECT */6 COLUMNS
SELECTION: [(col("comment_karma")) > (0)]
```

## The optimized plan is to:

- read the data from the Reddit CSV
- apply the filter on the comment\_karma column while the CSV is being read line-by-line
- transform the name column to uppercase

In this case the query optimizer has identified that the filter can be applied while the CSV is read from disk rather than reading the whole file into memory and then applying the filter. This optimization is called *Predicate Pushdown*.

## 4.6.5 Query execution

Our example query on the Reddit dataset is:



```
API scan_csv
```

```
q1 = (
    pl.scan_csv("docs/src/data/reddit.csv")
    .with_columns(pl.col("name").str.to_uppercase())
    .filter(pl.col("comment_karma") > 0)
)
```

If we were to run the code above on the Reddit CSV the query would not be evaluated. Instead Polars takes each line of code, adds it to the internal query graph and optimizes the query graph.

When we execute the code Polars executes the optimized query graph by default.

#### EXECUTION ON THE FULL DATASET

We can execute our query on the full dataset by calling the .collect method on the query.



#### API scan\_csv API collect

```
q4 = (
    pl.scan_csv(f"docs/data/reddit.csv")
    .with_columns(pl.col("name").str.to_uppercase())
    .filter(pl.col("comment_karma") > 0)
    .collect()
)
```

id	name	created utc	undated on	comment karma	link karma
		created_utt	updated_on	Comment_karma	IIIIK_Kalilla
					1
164	str	164	i64	i64	i64
<del></del>					
6	TAOJIANLONG_JASONBROKEN	1397113510	1536527864	4	0
17	SSAIG_JASONBROKEN	1397113544	1536527864	1	0
19	FDBVFDSSDGFDS JASONBROKEN	1397113552	1536527864	3	0
37	IHATEWHOWEARE JASONBROKEN	1397113636	1536527864	61	0
1229384	DSFOX	1163177415	1536497412	44411	7917
1229459	NEOCARTY	1163177859	1536533090	40	0
1229587	TEHSMA	1163178847	1536497412	14794	5707
1229621	JEREMYLOW	1163179075	1536497412	411	1063

Above we see that from the 10 million rows there are 14,029 rows that match our predicate.

With the default collect method Polars processes all of your data as one batch. This means that all the data has to fit into your available memory at the point of peak memory usage in your query.

#### **EXECUTION ON LARGER-THAN-MEMORY DATA**

If your data requires more memory than you have available Polars may be able to process the data in batches using *streaming* mode. To use streaming mode you simply pass the streaming=True argument to collect



```
API scan_csv API collect
```

```
q5 = (
    pl.scan_csv(f"docs/data/reddit.csv")
    .with_columns(pl.col("name").str.to_uppercase())
    .filter(pl.col("comment_karma") > 0)
    .collect(streaming=True)
)
```

We look at streaming in more detail here.

#### **EXECUTION ON A PARTIAL DATASET**

While you're writing, optimizing or checking your query on a large dataset, querying all available data may lead to a slow development process.

You can instead execute the query with the .fetch method. The .fetch method takes a parameter n\_rows and tries to 'fetch' that number of rows at the data source. The number of rows cannot be guaranteed, however, as the lazy API does not count how many rows there are at each stage of the query.

Here we "fetch" 100 rows from the source file and apply the predicates.



```
API scan_csv API collect API fetch
```

```
q9 = (
   pl.scan_csv(f"docs/data/reddit.csv")
   .with_columns(pl.col("name").str.to_uppercase())
   .filter(pl.col("comment_karma") > 0)
   .fetch(n_rows=int(100))
)
```

id	name	created_utc	updated_on	comment_karma	link_karma
i64	str	164	i64	164	i64
6	TAOJIANLONG JASONBROKEN	1397113510	1536527864	4	0
17	SSAIG JASONBROKEN	1397113544	1536527864	1	0
19	FDBVFDSSDGFDS_JASONBROKEN	1397113552	1536527864	3	0
37	IHATEWHOWEARE_JASONBROKEN	1397113636	1536527864	61	0
	l				
77763	LUNCHY	1137599510	1536528275	65	0
77765	COMPOSTELLAS	1137474000	1536528276	6	0
77766	GENERICBOB	1137474000	1536528276	291	14
77768	TINHEADNED	1139665457	1536497404	4434	103

# 4.6.6 Streaming



This section is still under development. Want to help out? Consider contributing and making a pull request to our repository. Please read our Contribution Guidelines on how to proceed.

# 4.7 IO

## 4.7.1 CSV

#### Read & Write

Reading a CSV file should look familiar:





API CsvReader Available on feature csv



API read\_csv

```
df = pl.read_csv("path.csv")
```

```
use polars::prelude::*;
let df = CsvReader::from_path("path.csv").unwrap().finish().unwrap();
```

API readCSV

```
df = pl.readCSV("path.csv")
```

Writing a CSV file is similar with the write\_csv function:







API write\_csv

```
df = pl.DataFrame({"foo": [1, 2, 3], "bar": [None, "bak", "baz"]})
df.write_csv("path.csv")
```

API CsvWriter · Available on feature csv

```
let mut df = df!(
   "foo" => &[1, 2, 3],
   "bar" => &[None, Some("bak"), Some("baz")],
)
.unwrap();
let mut file = std::fs::File::create("path.csv").unwrap();
CsvWriter::new(&mut file).finish(&mut df).unwrap();
```

API writeCSV

```
df = pl.DataFrame({ foo: [1, 2, 3], bar: [null, "bak", "baz"] });
df.writeCSV("path.csv");
```

#### Scan

Polars allows you to scan a CSV input. Scanning delays the actual parsing of the file and instead returns a lazy computation holder called a LazyFrame.



If you want to know why this is desirable, you can read more about these Polars optimizations here.

## 4.7.2 Parquet

Loading or writing Parquet files is lightning fast. Pandas uses PyArrow - Python bindings exposed by Arrow - to load Parquet files into memory, but it has to copy that data into Pandas memory. With Polars there is no extra cost due to copying as we read Parquet directly into Arrow memory and keep it there.

#### Read



## Write



## API write\_parquet

```
df = pl.DataFrame({"foo": [1, 2, 3], "bar": [None, "bak", "baz"]})
df.write_parquet("path.parquet")
```

API ParquetWriter · Available on feature parquet

```
let mut df = df!(
    "foo" => &[1, 2, 3],
    "bar" => &[None, Some("bak"), Some("baz")],
)
.unwrap();

let mut file = std::fs::File::create("path.parquet").unwrap();
ParquetWriter::new(&mut file).finish(&mut df).unwrap();
```

#### API writeParquet

```
df = pl.DataFrame({ foo: [1, 2, 3], bar: [null, "bak", "baz"] });
df.writeParquet("path.parquet");
```

### Scan

df = pl.scanParquet("path.parquet");

Polars allows you to scan a Parquet input. Scanning delays the actual parsing of the file and instead returns a lazy computation holder called a LazyFrame.



If you want to know why this is desirable, you can read more about those Polars optimizations here.

# 4.7.3 JSON files

### Read & Write

**JSON** 

Reading a JSON file should look familiar:





API read\_json

```
df = pl.read_json("path.json")
```

API JsonReader · Available on feature json

```
use polars::prelude::*;
let mut file = std::fs::File::open("path.json").unwrap();
let df = JsonReader::new(&mut file).finish().unwrap();
```

### NEWLINE DELIMITED JSON

JSON objects that are delimited by newlines can be read into polars in a much more performant way than standard json.





API read\_ndjson

```
df = pl.read_ndjson("path.json")
```

API JsonLineReader · Available on feature json

```
let mut file = std::fs::File::open("path.json").unwrap();
let df = JsonLineReader::new(&mut file).finish().unwrap();
```

### Write



API write\_json API write\_ndjson

```
df = pl.DataFrame({"foo": [1, 2, 3], "bar": [None, "bak", "baz"]})
df.write_json("path.json")
df.write_ndjson("path.json")
```

API JsonWriter API JsonWriter • Available on feature json

```
let mut df = df!(
    "foo" => &[1, 2, 3],
"bar" => &[None, Some("bak"), Some("baz")],
let mut file = std::fs::File::create("path.json").unwrap();
JsonWriter::new(&mut file)
   .with_json_format(JsonFormat::Json)
    .finish(&mut df)
    .unwrap();
// ndjson
JsonWriter::new(&mut file)
   .with_json_format(JsonFormat::JsonLines)
.finish(&mut df)
    .unwrap();
```

## Scan

Polars allows you to scan a JSON input only for newline delimited json. Scanning delays the actual parsing of the file and instead returns a lazy computation holder called a LazyFrame.





API scan\_ndjson

df = pl.scan\_ndjson("path.json")



let df = LazyJsonLineReader::new("path.json".to\_string()).finish().unwrap();

## 4.7.4 Multiple

## Dealing with multiple files.

Polars can deal with multiple files differently depending on your needs and memory strain.

Let's create some files to give us some context:

```
Python
```

API write\_csv

```
import polars as pl

df = pl.DataFrame({"foo": [1, 2, 3], "bar": [None, "ham", "spam"]})

for i in range(5):
    df.write_csv(f"my_many_files_{i}.csv")
```

## Reading into a single DataFrame

To read multiple files into a single  ${\tt DataFrame}$  , we can use globbing patterns:

```
Python
```

API read\_csv


```
df = pl.read_csv("my_many_files_*.csv")
print(df)
```

To see how this works we can take a look at the query plan. Below we see that all files are read separately and concatenated into a single DataFrame. Polars will try to parallelize the reading.



API show\_graph

```
pl.scan_csv("my_many_files_*.csv").show_graph()
```



# Reading and processing in parallel

If your files don't have to be in a single table you can also build a query plan for each file and execute them in parallel on the Polars thread pool.

All query plan execution is embarrassingly parallel and doesn't require any communication.



### API scan\_csv

```
import polars as pl
import glob

queries = []
for file in glob.glob("my_many_files_*.csv"):
    q = pl.scan_csv(file).groupby("bar").agg([pl.count(), pl.sum("foo")])
    queries.append(q)

dataframes = pl.collect_all(queries)
print(dataframes)
```

```
[shape: (3, 3)
 bar
                  foo
 str
         u32
                  i64
 null
 spam
                  2
 ham
                    ┙, shape: (3, 3)
 bar
                  foo
 str
         u32
                  i64
 null
                 2
         1
 spam
                     ┙, shape: (3, 3)
 bar
                  foo
 str
         u32
                  i64
 null
                  2
 ham
 spam
                        shape: (3, 3)
 bar
                  foo
 str
         u32
                  i64
 null
                 3
 spam
ham
                        shape: (3, 3)
 bar
                  foo
 str
         u32
                  i64
 ham
 spam
null
```

### 4.7.5 Databases

### Read from a database

We can read from a database with Polars using the pl.read\_database function. To use this function you need an SQL query string and a connection string called a connection\_uri.

For example, the following snippet shows the general patterns for reading all columns from the foo table in a Postgres database:



API read\_database · Available on feature connectorx

```
import polars as pl
connection_uri = "postgres://username:password@server:port/database"
query = "SELECT * FROM foo"
pl.read_database(query=query, connection_uri=connection_uri)
```

#### **ENGINES**

Polars doesn't manage connections and data transfer from databases by itself. Instead external libraries (known as *engines*) handle this. At present Polars can use two engines to read from databases:

- · ConnectorX and
- ADBC

### ConnectorX

ConnectorX is the default engine and supports numerous databases including Postgres, Mysql, SQL Server and Redshift. ConnectorX is written in Rust and stores data in Arrow format to allow for zero-copy to Polars.

To read from one of the supported databases with ConnectorX you need to activate the additional dependancy ConnectorX when installing Polars or install it manually with

```
$ pip install connectorx
```

### ADBC

ADBC (Arrow Database Connectivity) is an engine supported by the Apache Arrow project. ADBC aims to be both an API standard for connecting to databases and libraries implementing this standard in a range of languages.

It is still early days for ADBC so support for different databases is still limited. At present drivers for ADBC are only available for Postgres and SQLite. To install ADBC you need to install the driver for your database. For example to install the driver for SQLite you run

```
$ pip install adbc-driver-sqlite
```

As ADBC is not the default engine you must specify the engine as an argument to pl.read\_database



### API read database

```
connection_uri = "postgres://username:password@server:port/database"
query = "SELECT * FROM foo"
pl.read_database(query=query, connection_uri=connection_uri, engine="adbc")
```

### Write to a database

We can write to a database with Polars using the pl.write\_database function.

### **ENGINES**

As with reading from a database above Polars uses an engine to write to a database. The currently supported engines are:

- · SQLAlchemy and
- Arrow Database Connectivity (ADBC)

### SQLAlchemy

With the default engine SQLAlchemy you can write to any database supported by SQLAlchemy. To use this engine you need to install SQLAlchemy and Pandas

```
$ pip install SQLAlchemy pandas
```

In this example, we write the DataFrame to a table called records in the database



# API write\_database

```
connection_uri = "postgres://username:password@server:port/database"
df = pl.DataFrame({"foo": [1, 2, 3]})
df.write_database(table_name="records", connection_uri=connection_uri)
```

In the SQLAlchemy approach Polars converts the DataFrame to a Pandas DataFrame backed by PyArrow and then uses SQLAlchemy methods on a Pandas DataFrame to write to the database.

#### ADRO

As with reading from a database you can also use ADBC to write to a SQLite or Posgres database. As shown above you need to install the appropriate ADBC driver for your database.



### API write\_database

```
connection_uri = "postgres://username:password@server:port/database"
df = pl.DataFrame({"foo": [1, 2, 3]})

df.write_database(table_name="records", connection_uri=connection_uri, engine="adbc")
```

## 4.7.6 AWS



This section is still under development. Want to help out? Consider contributing and making a pull request to our repository. Please read our Contribution Guidelines on how to proceed.

To read from or write to an AWS bucket, additional dependencies are needed in Rust:



```
$ cargo add aws_sdk_s3 aws_config tokio --features tokio/full
```

In the next few snippets we'll demonstrate interacting with a Parquet file located on an AWS bucket.

#### Read

Load a .parquet file using:





API from\_arrow Available on feature fsspec Available on feature pyarrow



```
import polars as pl
import pyarrow.parquet as pq
import s3fs
fs = s3fs.S3FileSystem()
bucket = "<YOUR_BUCKET>"
path = "<YOUR_PATH>"
dataset = pq.ParquetDataset(f"s3://{bucket}/{path}", filesystem=fs)
df = pl.from_arrow(dataset.read())
```

```
use aws_sdk_s3::Region;
use aws_config::meta::region::RegionProviderChain;
use std::borrow::Cow;
use polars::prelude::*;
#[tokio::main]
async fn main() {
   let bucket = "<YOUR_BUCKET>";
   let path = "<YOUR_PATH>";
    let config = aws_config::from_env().load().await;
let client = Client::new(&config);
     let req = client.get_object().bucket(bucket).key(path);
     let res = req.clone().send().await.unwrap();
    let bytes = res.body.collect().await.unwrap();
let bytes = bytes.into_bytes();
     let cursor = std::io::Cursor::new(bytes);
     let df = CsvReader::new(cursor).finish().unwrap();
     println!("{:?}", df);
```

# 4.7.7 Google BigQuery

To read or write from GBQ, additional dependencies are needed:



\$ pip install google-cloud-bigquery

### Read

We can load a query into a DataFrame like this:





```
import polars as pl
from google.cloud import bigquery

client = bigquery.Client()

# Perform a query.
QUERY = (
    'SELECT name FROM 'bigquery-public-data.usa_names.usa_1910_2013` '
    'WHERE state = "TX" '
    'LIMIT 100')
query_job = client.query(QUERY)  # API request
rows = query_job.result()  # Waits for query to finish

df = pl.from_arrow(rows.to_arrow())
```

## Write



This section is still under development. Want to help out? Consider contributing and making a pull request to our repository. Please read our Contribution Guidelines on how to proceed.

# 4.8 SQL

### 4.8.1 Introduction

While Polars does support writing queries in SQL, it's recommended that users familiarize themselves with the expression syntax for more readable and expressive code. As a primarily DataFrame library, new features will typically be added to the expression API first. However, if you already have an existing SQL codebase or prefer to use SQL, Polars also offers support for SQL queries.



In Polars, there is no separate SQL engine because Polars translates SQL queries into expressions, which are then executed using its built-in execution engine. This approach ensures that Polars maintains its performance and scalability advantages as a native DataFrame library while still providing users with the ability to work with SQL queries.

#### Context

Polars uses the SQLContext to manage SQL queries . The context contains a dictionary mapping DataFrames and LazyFrames names to their corresponding datasets  $^1$ . The example below starts a SQLContext:



API SQLContext

```
ctx = pl.SQLContext()
```

### **Register Dataframes**

There are 2 ways to register DataFrames in the  ${\tt SQLContext}$ :

- register all LazyFrames and DataFrames in the global namespace
- register them one by one



# API SQLContext

```
df = pl.DataFrame({"a": [1, 2, 3]})
lf = pl.LazyFrame({"b": [4, 5, 6]})

# Register all dataframes in the global namespace: registers both df and lf
ctx = pl.SQLContext(register_globals=True)

# Other option: register dataframe df as "df" and lazyframe lf as "lf"
ctx = pl.SQLContext(df=df, lf=lf)
```

We can also register Pandas DataFrames by converting them to Polars first.



## API SQLContext

```
import pandas as pd

df_pandas = pd.DataFrame({"c": [7, 8, 9]})

ctx = pl.SQLContext(df_pandas=pl.from_pandas(df_pandas))
```



Converting a Pandas DataFrame backed by Numpy to Polars triggers a conversion to the Arrow format. This conversion has a computation cost. Converting a Pandas DataFrame backed by Arrow on the other hand will be free or almost free.

Once the SQLContext is initialized, we can register additional Dataframes or unregister existing Dataframes with:

- register
- register globals
- register\_many
- unregister

### **Execute queries and collect results**

SQL queries are always executed in lazy mode to benefit from lazy optimizations, so we have 2 options to collect the result:


- Set the parameter eager\_execution to True in SQLContext . With this parameter, Polars will automatically collect SQL results
- Set the parameter eager to True when executing a query with execute, or collect the result with collect.

We execute  $SQL\ queries\ by\ calling\ \ \mbox{execute}\ \ on\ a\ \ \mbox{SQLContext}\ .$ 



API register API execute

```
# For local files use scan_csv instead
pokemon = pl.read_csv(
    "https://gist.githubusercontent.com/ritchie46/cac6b337ea52281aa23c049250a4ff03/raw/89a957ff3919d90e6ef2d34235e6bf22304f3366/pokemon.csv"
)
ctx = pl.SQLContext(register_globals=True, eager_execution=True)
df_small = ctx.execute("SELECT * from pokemon LIMIT 5")
print(df_small)
```



### Execute queries from multiple sources

SQL queries can be executed just as easily from multiple sources. In the example below, we register:

- · a CSV file loaded lazily
- a NDJSON file loaded lazily
- a Pandas DataFrame

And we join them together with SQL. Lazy reading allows to only load the necessary rows and columns from the files.

In the same way, it's possible to register cloud datalakes (S3, Azure Data Lake). A PyArrow dataset can point to the datalake, then Polars can read it with scan pyarrow dataset.



### API register API execute

```
# Input data:
# products_masterdata.csv with schema {'product_id': Int64, 'product_name': Utf8}
# products_categories.json with schema ('product_id': Int64, 'category': Utf8)
# sales_data is a Pandas DataFrame with schema ('product_id': Int64, 'sales': Int64)

ctx = pl.SQLContext(
    products_masterdata=pl.scan_csv("products_masterdata.csv"),
    products_categories=pl.scan_ndjson("products_categories.json"),
    sales_data=pl.from_pandas(sales_data),
    eager_execution=True,
}

query = """

SELECT
    product_id,
    product_id,
    product_name,
    category,
    sales
FROM
    products_masterdata

LEFT_JOIN products_categories_USING (product_id)

LEFT_JOIN sales_data_USING (product_id)
"""

print(ctx.execute(query))
```

nape: (5, 4)			
product_id	product_name	category	sales
i64	str	str	i64
1	Product A	Category 1	100
2	Product B	Category 1	200
3	Product C	Category 2	150
4	Product D	Category 2	250
5	Product E	Category 3	300

### Compatibility

Polars does not support the full SQL language, in Polars you are allowed to:

- Write a CREATE statements CREATE TABLE XXX AS ...
- Write a SELECT statements with all generic elements ( GROUP BY , WHERE , ORDER , LIMIT , JOIN , ...)
- Write Common Table Expressions (CTE's) ( WITH tablename AS )
- Show an overview of all tables SHOW TABLES

The following is not yet supported:

- INSERT, UPDATE or DELETE statements
- $\bullet$  Table aliasing (e.g. SELECT p.Name from pokemon AS p)
- Meta queries such as ANALYZE, EXPLAIN

In the upcoming sections we will cover each of the statements in more details.

<sup>1.</sup> Additionally it also tracks the common table expressions as well.  $\leftarrow$ 

### 4.8.2 SHOW TABLES

In Polars, the Show Tables statement is used to list all the tables that have been registered in the current SQLContext. When you register a DataFrame with the SQLContext, you give it a name that can be used to refer to the DataFrame in subsequent SQL statements. The Show Tables statement allows you to see a list of all the registered tables, along with their names.

The syntax for the SHOW TABLES statement in Polars is as follows:

```
SHOW TABLES
```

Here's an example of how to use the SHOW TABLES statement in Polars:



```
API register API execute
```

In this example, we create two DataFrames and register them with the SQLContext using different names. We then execute a SHOW TABLES statement using the execute() method of the SQLContext object, which returns a DataFrame containing a list of all the registered tables and their names. The resulting DataFrame is then printed using the print() function.

Note that the SHOW TABLES statement only lists tables that have been registered with the current SQLContext. If you register a DataFrame with a different SQLContext or in a different Python session, it will not appear in the list of tables returned by SHOW TABLES.

## **4.8.3 SELECT**

In Polars SQL, the SELECT statement is used to retrieve data from a table into a DataFrame. The basic syntax of a SELECT statement in Polars SQL is as follows:

```
SELECT column1, column2, ...
FROM table_name;
```

Here, column1, column2, etc. are the columns that you want to select from the table. You can also use the wildcard \* to select all columns. table\_name is the name of the table or that you want to retrieve data from. In the sections below we will cover some of the more common SELECT variants



### API register API execute

```
shape: (6, 3)
 city
                              population
                              i64
New York
                              8399000
 Los Angeles
               USA
                              3997000
               USA
 Chicago
               USA
 Houston
 Phoenix
               USA
                              1680000
 Amsterdam
               Netherlands
                              900000
```

# GROUP BY

The GROUP BY statement is used to group rows in a table by one or more columns and compute aggregate functions on each group.



### API execute

```
result = ctx.execute(
    """
        SELECT country, AVG(population) as avg_population
        FROM population
        GROUP BY country
    """
)
print(result)
```

### ORDER BY

The ORDER BY statement is used to sort the result set of a query by one or more columns in ascending or descending order.



### API execute

```
result = ctx.execute(
    """
    SELECT city, population
    FROM population
    ORDER BY population
    """
)
print(result)
```

JOIN



### API register\_many API execute

```
income = pl.DataFrame(
         "city": [
"New York",
             "Los Angeles",
             "Chicago",
"Houston",
             "Amsterdam",
              "Rotterdam",
              "Utrecht",
          "country": [
"USA",
             "USA",
"USA",
"USA",
              "Netherlands",
              "Netherlands"
              "Netherlands",
         ],
"income": [55000, 62000, 48000, 52000, 42000, 38000, 41000],
ctx.register_many(income=income)
result = ctx.execute(
        SELECT country, city, income, population
        FROM population
LEFT JOIN income on population.city = income.city
print(result)
```

```
shape: (6, 4)
                                      population
               city
country
                             income
 str
               str
                             i64
                                       i64
 USA
                                       8399000
               New York
 USA
               Los Angeles
                             62000
                                       3997000
 USA
               Chicago
                             48000
 USA
                             52000
                                       2320000
               Houston
USA
              Phoenix
                                       1680000
```

```
| Netherlands | Amsterdam | 42000 | 900000 |
```

#### **FUNCTIONS**

Polars provides a wide range of SQL functions, including:

- $\bullet$  Mathematical functions: ABS , EXP , LOG , ASIN , ACOS , ATAN , etc.
- String functions: LOWER, UPPER, LTRIM, RTRIM, STARTS WITH, ENDS WITH.
- $\bullet$  Aggregation functions: SUM , AVG , MIN , MAX , COUNT , STDDEV , FIRST etc.
- $\bullet$  Array functions: <code>EXPLODE</code> , <code>UNNEST</code> , <code>ARRAY\_SUM</code> , <code>ARRAY\_REVERSE</code> , etc.

For a full list of supported functions go the API documentation. The example below demonstrates how to use a function in a query



### API query

```
result = ctx.execute(
    """

    SELECT city, population
    FROM population
    WHERE STARTS_WITH(country,'U')
    """
)
print(result)
```

```
city | population | --- | --- | str | i64 | | New York | 8399000 | Los Angeles | 3997000 | Chicago | 2705000 | Houston | 2320000 | Phoenix | 1680000 |
```

### TABLE FUNCTIONS

In the examples earlier we first generated a DataFrame which we registered in the SQLContext. Polars also support directly reading from CSV, Parquet, JSON and IPC in your SQL query using table functions read xxx.



# API execute

```
shape: (150, 5)
                                                                   species
 sepal_length
                  sepal_width
                                  petal_length
                                                   petal_width
 f64
                  f64
                                  f64
                                                   f64
                                                                   str
 5.1
                                                   0.2
                                                                   Setosa
 4.9
                  3.0
                                  1.4
                                                   0.2
                                                                   Setosa
                  3.2
                                                   0.2
 4.6
                  3.1
                                  1.5
                                                   0.2
                                                                   Setosa
 6.3
                  2.5
                                  5.0
                                                   1.9
                                                                   Virginica
                                                   2.0
                                                                   Virginica
Virginica
 6.5
                                  5.2
 6.2
                  3.4
                                  5.4
```

## **4.8.4 CREATE**

In Polars, the SQLContext provides a way to execute SQL statements against LazyFrames and DataFrames using SQL syntax. One of the SQL statements that can be executed using SQLContext is the CREATE TABLE statement, which is used to create a new table.

The syntax for the CREATE TABLE statement in Polars is as follows:

```
CREATE TABLE table_name
AS
SELECT ...
```

In this syntax, table\_name is the name of the new table that will be created, and SELECT ... is a SELECT statement that defines the data that will be inserted into the table.

Here's an example of how to use the CREATE TABLE statement in Polars:



### API register API execute

```
data = {"name": ["Alice", "Bob", "Charlie", "David"], "age": [25, 30, 35, 40]}
df = pl.LazyFrame(data)

ctx = pl.SQLContext(my_table=df, eager_execution=True)

result = ctx.execute(
    """
    CREATE TABLE older_people
    AS
    SELECT * FROM my_table WHERE age > 30

"""

print(ctx.execute("SELECT * FROM older_people"))
```

In this example, we use the <code>execute()</code> method of the <code>SQLContext</code> to execute a <code>CREATE TABLE</code> statement that creates a new table called <code>older\_people</code> based on a SELECT statement that selects all rows from the <code>my\_table</code> DataFrame where the <code>age column</code> is greater than 30.



Note that the result of a CREATE TABLE statement is not the table itself. The table is registered in the SQLContext . In case you want to turn the table back to a DataFrame you can use a SELECT  $\star$  FROM ... statement

## 4.8.5 Common Table Expressions

Common Table Expressions (CTEs) are a feature of SQL that allow you to define a temporary named result set that can be referenced within a SQL statement. CTEs provide a way to break down complex SQL queries into smaller, more manageable pieces, making them easier to read, write, and maintain.

A CTE is defined using the WITH keyword followed by a comma-separated list of subqueries, each of which defines a named result set that can be used in subsequent queries. The syntax for a CTE is as follows:

```
WITH cte_name AS (
subquery
)
SELECT ...
```

In this syntax, cte\_name is the name of the CTE, and subquery is the subquery that defines the result set. The CTE can then be referenced in subsequent queries as if it were a table or view.

CTEs are particularly useful when working with complex queries that involve multiple levels of subqueries, as they allow you to break down the query into smaller, more manageable pieces that are easier to understand and debug. Additionally, CTEs can help improve query performance by allowing the database to optimize and cache the results of subqueries, reducing the number of times they need to be executed.

Polars supports Common Table Expressions (CTEs) using the WITH clause in SQL syntax. Below is an example



### API register API execute

In this example, we use the <code>execute()</code> method of the <code>sqlContext</code> to execute a SQL query that includes a CTE. The CTE selects all rows from the <code>my\_table</code> LazyFrame where the <code>age</code> column is greater than 30 and gives it the alias <code>older\_people</code>. We then execute a second SQL query that selects all rows from the <code>older\_people</code> CTE where the <code>name</code> column starts with the letter 'C'.

# 4.9 Migrating

### 4.9.1 Coming from Pandas

Here we set out the key points that anyone who has experience with Pandas and wants to try Polars should know. We include both differences in the concepts the libraries are built on and differences in how you should write Polars code compared to Pandas code.

#### Differences in concepts between Polars and Pandas

POLARS DOES NOT HAVE A MULTI-INDEX/INDEX

Pandas gives a label to each row with an index. Polars does not use an index and each row is indexed by its integer position in the table.

Polars aims to have predictable results and readable queries, as such we think an index does not help us reach that objective. We believe the semantics of a query should not change by the state of an index or a reset index call.

In Polars a DataFrame will always be a 2D table with heterogeneous data-types. The data-types may have nesting, but the table itself will not. Operations like resampling will be done by specialized functions or methods that act like 'verbs' on a table explicitly stating the columns that that 'verb' operates on. As such, it is our conviction that not having indices make things simpler, more explicit, more readable and less error-prone.

Note that an 'index' data structure as known in databases will be used by polars as an optimization technique.

POLARS USES APACHE ARROW ARRAYS TO REPRESENT DATA IN MEMORY WHILE PANDAS USES NUMPY ARRAYS

Polars represents data in memory with Arrow arrays while Pandas represents data in memory with Numpy arrays. Apache Arrow is an emerging standard for inmemory columnar analytics that can accelerate data load times, reduce memory usage and accelerate calculations.

Polars can convert data to Numpy format with the to numpy method.

POLARS HAS MORE SUPPORT FOR PARALLEL OPERATIONS THAN PANDAS

Polars exploits the strong support for concurrency in Rust to run many operations in parallel. While some operations in Pandas are multi-threaded the core of the library is single-threaded and an additional library such as Dask must be used to parallelize operations.

POLARS CAN LAZILY EVALUATE QUERIES AND APPLY QUERY OPTIMIZATION

Eager evaluation is when code is evaluated as soon as you run the code. Lazy evaluation is when running a line of code means that the underlying logic is added to a query plan rather than being evaluated.

Polars supports eager evaluation and lazy evaluation whereas Pandas only supports eager evaluation. The lazy evaluation mode is powerful because Polars carries out automatic query optimization when it examines the query plan and looks for ways to accelerate the query or reduce memory usage.

Dask also supports lazy evaluation when it generates a query plan. However, Dask does not carry out query optimization on the query plan.

# Key syntax differences

Users coming from Pandas generally need to know one thing...

polars != pandas

If your Polars code looks like it could be Pandas code, it might run, but it likely runs slower than it should.

Let's go through some typical Pandas code and see how we might rewrite it in Polars.

SELECTING DATA

As there is no index in Polars there is no .loc or iloc method in Polars - and there is also no SettingWithCopyWarning in Polars.

However, the best way to select data in Polars is to use the expression API. For example, if you want to select a column in Pandas you can do one of the following:

```
df['a']
df.loc[:,'a']
```

but in Polars you would use the .select method:

```
df.select('a')
```

If you want to select rows based on the values then in Polars you use the .filter method:

```
df.filter(pl.col('a') < 10)</pre>
```

As noted in the section on expressions below, Polars can run operations in .select and filter in parallel and Polars can carry out query optimization on the full set of data selection criteria.

#### BE LAZY

Working in lazy evaluation mode is straightforward and should be your default in Polars as the lazy mode allows Polars to do query optimization.

We can run in lazy mode by either using an implicitly lazy function (such as scan\_csv) or explicitly using the lazy method.

Take the following simple example where we read a CSV file from disk and do a groupby. The CSV file has numerous columns but we just want to do a groupby on one of the id columns (idl) and then sum by a value column (v1). In Pandas this would be:

```
df = pd.read_csv(csv_file, usecols=['id1','v1'])
grouped_df = df.loc[:,['id1','v1']].groupby('id1').sum('v1')
```

In Polars you can build this query in lazy mode with query optimization and evaluate it by replacing the eager Pandas function read\_csv with the implicitly lazy Polars function scan\_csv:

```
df = pl.scan_csv(csv_file)
grouped_df = df.groupby('id1').agg(pl.col('v1').sum()).collect()
```

Polars optimizes this query by identifying that only the idl and v1 columns are relevant and so will only read these columns from the CSV. By calling the .collect method at the end of the second line we instruct Polars to eagerly evaluate the query.

If you do want to run this query in eager mode you can just replace scan\_csv with read\_csv in the Polars code.

Read more about working with lazy evaluation in the lazy API section.

### EXPRESS YOURSELF

A typical Pandas script consists of multiple data transformations that are executed sequentially. However, in Polars these transformations can be executed in parallel using expressions.

### Column assignment

We have a dataframe df with a column called value. We want to add two new columns, a column called tenXValue where the value column is multiplied by 10 and a column called hundredXValue where the value column is multiplied by 100.

In Pandas this would be:

```
df["tenXValue"] = df["value"] * 10
df["hundredXValue"] = df["value"] * 100
```

These column assignments are executed sequentially.

In Polars we add columns to df using the .with\_columns method and name them with the .alias method:

```
df.with_columns(
    (pl.col("value") * 10).alias("tenXValue"),
    (pl.col("value") * 100).alias("hundredXValue"),
)
```

These column assignments are executed in parallel.

#### Column assignment based on predicate

In this case we have a dataframe df with columns a, b and c. We want to re-assign the values in column a based on a condition. When the value in column c is equal to 2 then we replace the value in a with the value in b.

In Pandas this would be:

```
df.loc[df["c"] == 2, "a"] = df.loc[df["c"] == 2, "b"]
```

while in Polars this would be:

```
df.with_columns(
    pl.when(pl.col("c") == 2)
    .then(pl.col("b"))
    .otherwise(pl.col("a")).alias("a")
)
```

The Polars way is pure in that the original DataFrame is not modified. The mask is also not computed twice as in Pandas (you could prevent this in Pandas, but that would require setting a temporary variable).

Additionally Polars can compute every branch of an if -> then -> otherwise in parallel. This is valuable, when the branches get more expensive to compute.

#### Filtering

We want to filter the dataframe df with housing data based on some criteria.

In Pandas you filter the dataframe by passing Boolean expressions to the loc method:

```
df.loc[(df['sqft_living'] > 2500) & (df['price'] < 300000)]
```

while in Polars you call the filter method:

```
df.filter(
     (pl.col("m2_living") > 2500) & (pl.col("price") < 300000)
)</pre>
```

The query optimizer in Polars can also detect if you write multiple filters separately and combine them into a single filter in the optimized plan.

### Pandas transform

The Pandas documentation demonstrates an operation on a groupby called transform. In this case we have a dataframe df and we want a new column showing the number of rows in each group.

In Pandas we have:

```
df = pd.DataFrame({
    "type": ["m", "n", "o", "m", "n", "n"],
    "c": [1, 1, 1, 2, 2, 2, 2],
})
df["size"] = df.groupby("c")["type"].transform(len)
```

Here Pandas does a groupby on "c", takes column "type", computes the group length and then joins the result back to the original DataFrame producing:

```
c type size
0 1 m 3
1 1 n 3
2 1 o 3
3 2 m 4
4 2 m 4
5 2 n 4
6 2 n 4
```


In Polars the same can be achieved with window functions:

```
df.select(
   pl.all(),
   pl.col("type").count().over("c").alias("size")
)
```

Because we can store the whole operation in a single expression, we can combine several window functions and even combine different groups!

Polars will cache window expressions that are applied over the same group, so storing them in a single select is both convenient **and** optimal. In the following example we look at a case where we are calculating group statistics over "c" twice:

```
df.select(
   pl.all(),
   pl.col("c").count().over("c").alias("size"),
   pl.col("c").sum().over("type").alias("sum"),
   pl.col("c").reverse().over("c").flatten().alias("reverse_type")
)
```



# Missing data

Pandas uses NaN and/or None values to indicate missing values depending on the dtype of the column. In addition the behaviour in Pandas varies depending on whether the default dtypes or optional nullable arrays are used. In Polars missing data corresponds to a null value for all data types.

For float columns Polars permits the use of NaN values. These NaN values are not considered to be missing data but instead a special floating point value.

In Pandas an integer column with missing values is cast to be a float column with NaN values for the missing values (unless using optional nullable integer dtypes). In Polars any missing values in an integer column are simply null values and the column remains an integer column.

See the missing data section for more details.

# 4.9.2 Coming from Apache Spark

### Column-based API vs. Row-based API

Whereas the Spark DataFrame is analogous to a collection of rows, a Polars DataFrame is closer to a collection of columns. This means that you can combine columns in Polars in ways that are not possible in Spark, because Spark preserves the relationship of the data in each row.

Consider this sample dataset:

```
import polars as pl

df = pl.DataFrame({
    "foo": ["a", "b", "c", "d", "d"],
    "bar": [1, 2, 3, 4, 5],
})

dfs = spark.createDataFrame(
    ("a", 1),
     ("b", 2),
     ("c", 3),
     ("d", 4),
     ("d", 5),
    ],
    schema=["foo", "bar"],
}
```

### EXAMPLE 1: COMBINING HEAD AND SUM

In Polars you can write something like this:

```
df.select(
    pl.col("foo").sort().head(2),
    pl.col("bar").filter(pl.col("foo") == "d").sum()
)
```

## Output:

The expressions on columns foo and bar are completely independent. Since the expression on bar returns a single value, that value is repeated for each value output by the expression on foo. But a and b have no relation to the data that produced the sum of 9.

To do something similar in Spark, you'd need to compute the sum separately and provide it as a literal:

```
from pyspark.sql.functions import col, sum, lit

bar_sum = (
    dfs
    .where(col("foo") == "d")
    .groupBy()
    .agg(sum(col("bar")))
    .take(1)[0][0]
)

(
    dfs
    .orderBy("foo")
    .limit(2)
    .withColumn("bar", lit(bar_sum))
    .show()
)
```

### Output:

```
+---+---+
|foo|bar|
+---+---+
| a| 9|
```

```
| b| 9|
+---+
```

### EXAMPLE 2: COMBINING TWO HEAD S

In Polars you can combine two different head expressions on the same DataFrame, provided that they return the same number of values.

```
df.select(
    pl.col("foo").sort().head(2),
    pl.col("bar").sort(descending=True).head(2),
)
```

### Output:

```
shape: (3, 2)

| foo | bar |
|--- |
|--- |
| str | i64 |
| a | 5 |
| b | 4 |
```

Again, the two head expressions here are completely independent, and the pairing of a to 5 and b to 4 results purely from the juxtaposition of the two columns output by the expressions.

To accomplish something similar in Spark, you would need to generate an artificial key that enables you to join the values in this way.

```
from pyspark.sql import Window
from pyspark.sql.functions import row_number

foo_dfs = (
    dfs
    .withColumn(
        "rownum",
        row_number().over(Window.orderBy("foo"))
    )
)

bar_dfs = (
    dfs
    .withColumn(
        "rownum",
        row_number().over(Window.orderBy(col("bar").desc()))
    )
)

(
    foo_dfs.alias("foo")
    .join(bar_dfs.alias("bar"), on="rownum")
    .select("foo.foo", "bar.bar")
    .limit(2)
    .show()
}
```

### Output:

```
+---+---+
|foolbar|
+---+---+
| a| 5|
| b| 4|
+---+---+
```

### 4.10 Misc

### 4.10.1 Multiprocessing

TLDR: if you find that using Python's built-in multiprocessing module together with Polars results in a Polars error about multiprocessing methods, you should make sure you are using <code>spawn</code>, not <code>fork</code>, as the starting method:

```
Python
```

```
from multiprocessing import get_context

def my_fun(s):
    print(s)

with get_context("spawn").Pool() as pool:
    pool.map(my_fun, ["input1", "input2", ...])
```

### When not to use multiprocessing

Before we dive into the details, it is important to emphasize that Polars has been built from the start to use all your CPU cores. It does this by executing computations which can be done in parallel in separate threads. For example, requesting two expressions in a <code>select</code> statement can be done in parallel, with the results only being combined at the end. Another example is aggregating a value within groups using <code>groupby().agg(<expr>)</code>, each group can be evaluated separately. It is very unlikely that the <code>multiprocessing</code> module can improve your code performance in these cases.

See the optimizations section for more optimizations.

### When to use multiprocessing

Although Polars is multithreaded, other libraries may be single-threaded. When the other library is the bottleneck, and the problem at hand is parallelizable, it makes sense to use multiprocessing to gain a speed up.

# The problem with the default multiprocessing config

SUMMARY

The Python multiprocessing documentation lists the three methods to create a process pool:

- 1. spawn
- 2. fork
- 3. forkserver

The description of fork is (as of 2022-10-15):

The parent process uses os.fork() to fork the Python interpreter. The child process, when it begins, is effectively identical to the parent process. All resources of the parent are inherited by the child process. Note that safely forking a multithreaded process is problematic.

Available on Unix only. The default on Unix.

The short summary is: Polars is multithreaded as to provide strong performance out-of-the-box. Thus, it cannot be combined with fork. If you are on Unix (Linux, BSD, etc), you are using fork, unless you explicitly override it.

The reason you may not have encountered this before is that pure Python code, and most Python libraries, are (mostly) single threaded. Alternatively, you are on Windows or MacOS, on which fork is not even available as a method (for MacOS it was up to Python 3.7).

Thus one should use spawn , or forkserver , instead. spawn is available on all platforms and the safest choice, and hence the recommended method.

### EXAMPLE

The problem with fork is in the copying of the parent's process. Consider the example below, which is a slightly modified example posted on the Polars issue tracker:



```
import multiprocessing
import polars as pl
def test_sub_process(df: pl.DataFrame, job_id):
    df_filtered = df.filter(pl.col("a") > 0)
    print(f"Filtered (job_id: {job_id})", df_filtered, sep="\n")
def create dataset():
    return pl.DataFrame({"a": [0, 2, 3, 4, 5], "b": [0, 4, 5, 56, 4]})
def setup():
     # some setup work
    df.write_parquet("/tmp/test.parquet")
def main():
    test_df = pl.read_parquet("/tmp/test.parquet")
    for i in range (0.5):
        proc = multiprocessing.get_context("spawn").Process(
            target=test_sub_process, args=(test_df, i)
        proc.start()
        proc.join()
        print(f"Executed sub process {i}")
if __name__ == "__main__":
    setup()
    main()
```

Using fork as the method, instead of spawn, will cause a dead lock. Please note: Polars will not even start and raise the error on multiprocessing method being set wrong, but if the check had not been there, the deadlock would exist.

The fork method is equivalent to calling os.fork(), which is a system call as defined in the POSIX standard:

A process shall be created with a single thread. If a multi-threaded process calls fork(), the new process shall contain a replica of the calling thread and its entire address space, possibly including the states of mutexes and other resources. Consequently, to avoid errors, the child process may only execute async-signal-safe operations until such time as one of the exec functions is called.

In contrast, spawn will create a completely new fresh Python interpreter, and not inherit the state of mutexes.

So what happens in the code example? For reading the file with pl.read\_parquet the file has to be locked. Then os.fork() is called, copying the state of the parent process, including mutexes. Thus all child processes will copy the file lock in an acquired state, leaving them hanging indefinitely waiting for the file lock to be released, which never happens.

What makes debugging these issues tricky is that fork can work. Change the example to not having the call to pl.read\_parquet:



This works fine. Therefore debugging these issues in larger code bases, i.e. not the small toy examples here, can be a real pain, as a seemingly unrelated change can break your multiprocessing code. In general, one should therefore never use the fork start method with multithreaded libraries unless there are very specific requirements that cannot be met otherwise.

PRO'S AND CONS OF FORK

Based on the example, you may think, why is fork available in Python to start with?

First, probably because of historical reasons: spawn was added to Python in version 3.4, whilst fork has been part of Python from the 2.x series.

Second, there are several limitations for spawn and forkserver that do not apply to fork, in particular all arguments should be pickable. See the Python multiprocessing does for more information.

Third, because it is faster to create new processes compared to spawn, as spawn is effectively fork + creating a brand new Python process without the locks by calling execv. Hence the warning in the Python docs that it is slower: there is more overhead to spawn. However, in almost all cases, one would like to use multiple processes to speed up computations that take multiple minutes or even hours, meaning the overhead is negligible in the grand scheme of things. And more importantly, it actually works in combination with multithreaded libraries.

Fourth, spawn starts a new process, and therefore it requires code to be importable, in contrast to fork. In particular, this means that when using spawn the relevant code should not be in the global scope, such as in Jupyter notebooks or in plain scripts. Hence in the examples above, we define functions where we spawn within, and run those functions from a main clause. This is not an issue for typical projects, but during quick experimentation in notebooks it could fail.

### References

- $1.\ https://docs.python.org/3/library/multiprocessing.html$
- $2.\ https://pythonspeed.com/articles/python-multiprocessing/$
- 3. https://pubs.opengroup.org/onlinepubs/9699919799/functions/fork.html
- 4. https://bnikolic.co.uk/blog/python/parallelism/2019/11/13/python-forkserver-preload.html

### 4.10.2 Alternatives

These are some tools that share similar functionality to what polars does.

#### Pandas

A very versatile tool for small data. Read 10 things I hate about pandas written by the author himself. Polars has solved all those 10 things. Polars is a versatile tool for small and large data with a more predictable, less ambiguous, and stricter API.

#### · Pandas the API

The API of pandas was designed for in memory data. This makes it a poor fit for performant analysis on large data (read anything that does not fit into RAM). Any tool that tries to distribute that API will likely have a suboptimal query plan compared to plans that follow from a declarative API like SQL or Polars' API.

#### • Dask

Parallelizes existing single-threaded libraries like NumPy and Pandas. As a consumer of those libraries Dask therefore has less control over low level performance and semantics. Those libraries are treated like a black box. On a single machine the parallelization effort can also be seriously stalled by pandas strings. Pandas strings, by default, are stored as python objects in numpy arrays meaning that any operation on them is GIL bound and therefore single threaded. This can be circumvented by multi-processing but has a non-trivial cost.

### • Modin

Similar to Dask

#### • Vaex

Vaexs method of out-of-core analysis is memory mapping files. This works until it doesn't. For instance parquet or csv files first need to be read and converted to a file format that can be memory mapped. Another downside is that the OS determines when pages will be swapped. Operations that need a full data shuffle, such as sorts, have terrible performance on memory mapped data. Polars' out of core processing is not based on memory mapping, but on streaming data in batches (and spilling to disk if needed), we control which data must be hold in memory, not the OS, meaning that we don't have unexpected IO stalls.

#### · DuckDB

Polars and DuckDB have many similarities. DuckDB is focused on providing an in-process OLAP Sqlite alternative, Polars is focused on providing a scalable DataFrame interface to many languages. Those different front-ends lead to different optimization strategies and different algorithm prioritization. The interoperability between both is zero-copy. See more: https://duckdb.org/docs/guides/python/polars

## • Spark

Spark is designed for distributed workloads and uses the JVM. The setup for spark is complicated and the startup-time is slow. On a single machine Polars has much better performance characteristics. If you need to process TB's of data Spark is a better choice.

### • CuDF

GPU's and CuDF are fast! However, GPU's are not readily available and expensive in production. The amount of memory available on a GPU is often a fraction of the available RAM. This (and out-of-core) processing means that Polars can handle much larger data-sets. Next to that Polars can be close in performance to CuDF. CuDF doesn't optimize your query, so is not uncommon that on ETL jobs Polars will be faster because it can elide unneeded work and materializations.

## • Any

Polars is written in Rust. This gives it strong safety, performance and concurrency guarantees. Polars is written in a modular manner. Parts of polars can be used in other query programs and can be added as a library.

# 4.10.3 Reference Guides

The api documentations with details on function / object signatures can be found here:

- NodeJS
- Python
- Rust

# 4.10.4 Contributing

See the CONTRIBUTING.md if you would like to contribute to the Polars project.

If you're new to this we recommend starting out with contributing examples to the Python API documentation. The Python API docs are generated from the docstrings of the Python wrapper located in polars/py-polars.

Here is an example commit that adds a docstring.

If you spot any gaps in this User Guide you can submit fixes to the pola-rs/polars-book repo.

Happy hunting!